

Deployments Playbook

SECOND EDITION

A guide to delivering virtual
assistants at PolyAI, for the
deployments team

V2.0.1
3 DEC 2025



Introduction

PolyAI builds virtual assistants, usually for voice, that empowers callers to engage with a technology that has a bad reputation. At their best, our assistants are indistinguishable from humans and drastically reduces the need for humans in a contact centre, reserving them for the callers that need empathy and tact.

Your job is to deliver these agents, to maximise the quality of them and to be client centric.

Our mission is to make everyday customer service feel superhuman, with voice assistants people actually like.

Your mission as the PSM is to turn this into a reality.

You will make our clients feel excited about the work they are doing. You'll chase ROI for our customers, and even when the measure is blurry, you'll chase engagement for their callers.

This playbook is your guide.

Introduction

This document will use icons and links to take you to further information.



You can click the link icon to find supplementary information, sample documents and code snippets.



This icon signals that this page is something that might soon change, for example, the product is evolving. Check with the team.

PART ONE

Project Delivery

What We Build

Voice Assistants that are Superhuman

PolyAI builds voice assistants. A good voice assistant is indistinguishable from a human for most tasks. The voices are respectful, natural, casual and convincing. They don't sound like robots, adverts, or technology.

We do offer text virtual assistants, in chat on the web, via SMS or WhatsApp, or in app. However we build *voice first* — this is because it's extremely hard to convert a chatbot to voice, but easier the other way round. Voice has many nuances such as background noise, crosstalk and bad signal. We therefore do not have any major live chat deployments today.

Voice Assistants that are Revenue Generating and Cost Saving

Our mission is not to replace contact centre agents, but to ensure that callers get through to humans only when they require additional empathy, care and tact. Contact centre agents should take the most interesting and complex calls. They also get the chance to work with us on their virtual agent.

We also can generate revenue. We are careful about how we measure this, as it's often a little out of our control. But we measure how many bookings we make, abandons we reduce, and try to link hold time reductions to revenue generation.

Generative vs Intent Based

Generative

- ✓ A generative AI system answers questions using a knowledge base that customers can edit
- ✓ Conversational flows can be created in PolyAI's Agent Studio, again with an LLM at its foundation
- ✓ Complex flows and high sensitivity ones like 'authentication and verification' flows may be built on PolyAI's prescriptive NLU model and policy framework by PolyAI engineers, but generally not.

Intent Based

- ✓ PolyAI takes a user input and maps it to an 'intent'. The system has prescribed behaviour, so takes a defined action based on that intent, which might be reading out an FAQ, or proceeding through a flow
- ✓ The answers the virtual agent gives are pre-written and fixed



You should be careful how you talk to clients about this. The exact mechanisms of how we build are not important. **PolyAI uses a hybrid system of both proprietary models and large language models, which are used interchangeably based on the customer's request.** Our systems are predictable and custom.

How We Build

We Build on a Generative AI Platform, called Agent Studio

Agent Studio is our platform and control panel for virtual assistants. It contains the knowledge base, flow design, voice generation, analytics, conversation review and more. Each of these different items can be enabled on a per-client basis.

Our platform has...

A configurable **Large Language Model**, which may be GPT, hosted on Azure AI, or our in-house LLM, known as Raven. We choose which to use on a per use case basis. Regardless, many client's will understand this as the PolyAI LLM.

A **PolyAI Speech Recognition System**. Designed and trained in house, this translates audio of human speech into text, which gets sent to the LLM. It's based on a base model from NVIDIA, who have invested in us, and we build on top of it.

A **lifelike hybrid Text to Speech model**. We offer samples to our clients that might be a mix of a PolyAI trained voice, based on an actor that we work closely with, off the shelf voices, or live recordings. Our client's don't pick, they trust us to use the best hybrid of all three options.

Who Builds


The Deployment Team is a subset of Team Y, a.k.a. The Project Team or the PS Org

It contains Agent Designers, Product Solutions Managers and the Support team.

Product Solutions Managers are the primary client contact for the build of a virtual agent. They own the timeline and the project plans. They build and manage dashboards, analytics, and success rates. They make the client feel good about what they've purchased.

Agent Designers script the agent and prompt the LLM. They are responsible for the quality of the agent and making it as lifelike as possible. They own relationships with the voice actors and produce deep analytics.

The Support team manage our service desk, where live and running clients submit fault requests.



Who Builds

Team Y, the wider Project Team, operates in 'pods'

Pods contain scrum masters and engineers, who report into Eddy's org, agent designers, PSMs and a CSM.

New projects and clients are assigned to the pods.

The pod makeups may change, but at time of writing are:

Bao, primarily based in the UK and responsible largely for FedEx

Sourdough, primarily based in the UK and responsible for many European clients

Somun, primarily Serbian engineers but responsible for US projects due to EU/UK data transfer rules

Focaccia, based in the UK with global projects


Bagel, based in the US and responsible for US projects and many casino projects

Paratha, based in both the UK and the US, and responsible for global projects



Crouton, based in west coast US, and responsible for PG&E and CLI

Deployment Process

This section will run through how to run a project, step by step, using a master diagram. You can view the large diagram by clicking below, or follow the process through the next few pages.



Deployment Process: Handovers



Sales to Delivery Handoff

This meeting is booked by the AE within 5-10 days of signature. At this point, a PSM is assigned to attend the meeting. Department heads will also join. This meeting has a structured slide deck. PSM's should ensure they know:


- Who the main points of contact are
- The scope, in detail, including integrations, parts that are missing, etc
- Expected timelines
- The commercial structure, ramp period, opt outs, etc

The External Kickoff Meeting

This meeting is booked by the PSM after the AE sends an introduction email. Usually, a client will invite all stakeholders to this meeting, as they expect to know who is responsible for each step of the process.

Details on how to run this meeting are found on page 28.

Deployment Process: Scope ①




Scope Run Through

After kickoff, you should validate the scope of the virtual agent with the client. Often, the contracted scope will have nuances that will change. We want to be easy to work with, so some changes should be accepted for free.


For this meeting you should:

- Be extremely familiar with the flow in the order form. A detail oriented PSM may re-create the diagram in a more client friendly way, slightly simplified.
- Not allow the client to leave too much in the open, you want answers, and if the client cannot answer them, this may delay build
- Be able to add sample API requests and responses, and after the meeting validate all of them in the diagram
- Get final written sign off over email that the scope is final.

This is a technical call and you should feel confident that what you create on paper can go to our engineering team, and be understood.




Deployment Process: Scope ②





Major Scope Changes

There may be instances where clients want to change the scope substantially. If the change would require any timeline change, or significantly more effort, there are two options.

1. With the CSM, draft a Scope Change Email , outlining the change. This is a formal scope change that the client must confirm. It may come at extra cost, but this is optional. It may also invalidate opt-out, or cause a timeline extension. Because of this, you should also check with legal before sending.
2. Escalate through the AE. If the scope of the project is entirely different to what was contracted, the project needs to be transferred back to the solutions consultant team, through the AE. Send a message on the commercial slack channel, and organise a meeting with relevant team members and your line manager.



Deployment Process: Tech Kickoff



Project Setup Checklist, For the PSM

- Create a Notion project page in the projects database. Use the template
- Input contracted milestone dates into the Notion
- Place links to all discovery and contractual information into the Notion
- Create a folder for Smartsheet items in the Project Timelines workspace
- Work with the scrum master to set up a Jira board and workflow
- Create _tech and _telephony slack channels



Compliance

It is possible there will be surprises for our compliance team, especially if the project collects PII (Personal Identifiable Information) or PHI (Personal Healthcare Information). It is often good to check with our compliance team, especially if it's a new industry or use case

Internal Tech Kickoff

By this point, a pod should be fully assigned, and the PSM organises an internal tech kickoff where they run through the agreed scope, and agreeing milestone dates. At this point, the engineering lead should begin planning how long each action might take, and working backwards from the contracted go live date.

Deployment Process: Timeline



Create the Client Facing Timeline

You will create a client facing timeline, and get sign off from engineering. There is a template available in Smartsheet  to use as a basis. It is important that this is up to date, day by day, and used to show the client when and how we are on track. The most effective way to hold a client to account is to assign them actions on this timeline.

You should add line items for every detailed action. You can use the lucid chart for this section of the playbook to help map out each line item. Be aware you should be able to explain them.

Work with the lead engineer in the pod on the dates.

Tasks							Timeline (Jul 25 to Oct 1)											
Task ID	Task Name	Assigned To	Start Date	End Date	Health	Status	Notes	Jul 25	Aug 1	Aug 8	Aug 15	Aug 22	Aug 29	Sep 5	Sep 12	Sep 19	Sep 26	Oct 1
1	Design and Scoping	PSM + AG	01/08/22	09/09/22	●	Not Started												
2	Project Kickoff Call		01/08/22	05/08/22	●	Not Started	To organise											
3	MI Discussion		01/08/22	05/08/22	●	Not Started	To discuss if aligned and											
4	Telephony Scoping	PSM	08/08/22	09/09/22	●	Not Started	PSTN - AG v PolyAI to pre											
5	Collect Call Routes, FAQs, Project Scope	PSM	08/08/22	19/09/22	●	Not Started	Nathan has											
6	Final Scope Signoff				●	Not Started	Working call											
7	Agent Design	DD	22/08/22	08/09/22	●	Not Started												
8	Content and Scope Review	PSM	09/09/22	09/09/22	●	Not Started	Prefer prese logic											
9	Voice Actor Reachout and Sign Off	DD	29/08/22	02/09/22	●	Not Started												
10	Agent Development		12/09/22	07/10/22														
11	Core Build		12/09/22	30/09/22		Not Started												
12	Policy	SM	12/09/22	30/09/22	●	Not Started												
13	Telephony Integration	SM	12/09/22	30/09/22	●	Not Started												
14	Metrics	SM + PSM	28/09/22	30/09/22	●	Not Started												
15	Initial Agent Complete	SM	30/09/22	30/09/22	●	Not Started												
16	Fine Tuning Period	SM	03/10/22	07/10/22	●	Not Started												
17	Voice Actor Recording	DD	19/09/22	23/09/22	●	Not Started												
18	Voice Actor Processing and Upload	DD	28/09/22	30/09/22	●	Not Started												
19	Build Dashboards	PSM	03/10/22	07/10/22	●	Not Started												
20	HLD and Go Live Checklist	SM	09/10/22	07/10/22	●	Not Started												
21	Testing		09/10/22	14/10/22														
22	UAT Test Plan	PSM	09/10/22	07/10/22	●	Not Started												
23	Client UAT	PSM	10/10/22	14/10/22	●	Not Started												
24	Final Client Sign Off	PSM	14/10/22	14/10/22	●	Not Started												
25	Rollout		17/10/22	18/11/22														
26	Enable System (Staggered Rollout)	PSM	17/10/22	21/10/22	●	Not Started												
27	Hypercare Support	PSM	24/10/22	18/11/22	●	Not Started												
28	Success Criteria Met	PSM	18/11/22	18/11/22	●	Not Started												



Using Smartsheet

Smartsheet can be tricky to use. However, as these are client facing it is imperative that they are clean, tidy, and professional. You will want to use this to hold clients to account, and if they are untidy or confusing, they will not be taken seriously.

Use the table view to enter dates, and then visualise it in the Gantt view afterwards. Do not set dependencies until you have finished setting dates, however do use the dependency feature to demonstrate to clients the cascading effect of them being late. Keep the status up to date at all times.

Use the Smartsheet [help guides](#).

Deployment Process: Building ①



Client Meetings



We run weekly meetings with the client, often more, to keep them up to date with progress. In these meetings you will show the action tracker, the timeline and report progress. You will play voice samples to the client and coach them into picking the right one. Work closely with the Agent Designer to pick the right voice.

We don't usually ask for preferences, rather send voices that we think would work for the agent, and then take feedback.

It's also often wise to get an early recording of the real system mid-development to show progress.

More guides on how to run such daily meetings can be found on page X.

Deployment Process: Building ①



Telephony Integration and Interoperability document

Your telephony engineer will own this document, however check in frequently to ensure it is up to date, and that the client is also following their deliverables on time. It is very easy for the deliverables in this document to slip.



Providing voice samples

Your Agent Designer will generate samples. Listen to them first, and make sure they are perfect. It's a little tedious to generate samples so don't be afraid to push back if you find the voice strange.

You should provide **no more than three** samples to the client upfront, and play them on a call (there is a template here [🔗](#)). The more you provide up front, the more limited the options if they don't like any.

You should always include 'PolyVoices' in the samples. These are voices where we own the relationship with the voice actor and so can get live recordings later for unusual pronunciations, or to maximise agent voice quality. Our clients have been sold superhuman assistants, so PolyVoices give us the most flexibility.

Deployment Process: Building ②





Keeping Close to the Engineering Team

WAs a PSM, the engineering team are ultimately waiting on your guidance and timelines. Projects will not move at the pace required unless you are close to the day to day. You must:

- Be deeply familiar with the bot and it's behaviour. Clients will ask you about it, and you should be able to answer forcefully and confidently. The bot's design is as much yours as it is the team's.
- Attend scrum to be aware of day to day deliverables
- Be present in the tech channel, if you are not seeing progress, there may not be progress. Nobody wants to micromanage, but you must be confident and having the team report progress. Delays, poor agent quality and client frustration will come if you cannot justify it.

Deployment Process: Metrics



Defining Metrics

The PSM is responsible for defining the metrics required in the project. There are some standard metrics which must be followed and you should check that these are implemented correctly.

Metrics are 'checkpoints' in a conversation. When a caller hits a point in a conversation, a metric will be written. It could be something simple like 'BOOKING_INITIATED' which is a single event. Or it could contain data, for example 'QA' which signals a question was asked but will have a value such as 'PARKING', the name of the FAQ.



Detail on standardised metrics can be found on Notion. [🔗](#) However, there is also a new effort to improve it, which will be documented soon. [🔄](#)

The Consequences of Incorrect Metrics

Metrics are how we measure business cases, and more, with our clients. Getting it wrong can create huge problems and require time consuming backfilling and extended manual call listening.

You should extensively QA that metrics are being written. QA documents must have an 'expected metrics' column, and you must validate these on every test call.

Deployment Process: QA and Review



Setting Up QA

The Agent Designer will create the initial QA sheet. This should be detailed and it should be considered client facing because the UAT sheet will be built from it.

It should cover every reasonable scenario, including:



- Happy paths
- API integrations, that should be mocked if not possible early in QA
- Unhappy paths and failures
- Standard behaviour such as 'speak to' and 'who are you'
- That all metrics are being written correctly

There is a template available in Agent Design Notion. A good model to check is that of Pradco / Moultrie Mobile.

Be Ready for Implementation Review

The implementation review board is a new process that will check the agent quality / completeness of QA / metrics and more. See page X for more detail. If the project is rejected at this stage, it may result in delaying UAT.

Deployment Process: UAT ①



Before UAT

Remember, the client is not here to do QA. Be ready, expect 100% pass rate from the team and expect no less. Prepare the QA sheet for client facing, with scoring using the below criteria. It is your responsibility to make the internal QA sheet client-friendly for UAT. See an example here. [🔗](#)

UAT Kickoff

Use the UAT slide template to explain the difference between functional and behavioural testing and plan testing. Reassure the client some things won't work but that's normal. You'll be making many changes over the period to improve the system. We'll listen to every call but priorities the bad CX fails.

Scoring Functional Testing

This is how to score a functional test you complete.

As Designed - Perfect

The outcome of the call was as designed. For example, a visitor who should be granted access was granted access. The user was understood first time, and the CX was good.

As Designed - Imperfect


The outcome of the call was as designed. For example, a visitor who should be granted access was granted access. However, the experience was not perfect. The outcome was correct, but the user may have had to repeat themselves.

Fail - Good UX

The outcome of the call was not as intended. For example, the agent transferred a user to a human operator despite the fact the visitor was authorized to enter and the API provided correct information. However, the user experienced a seamless experience, such as a fast transfer to a human.

Fail - Bad UX or Incorrect Admission

The bot did not behave as intended, and the experience was bad, for example, the user was stuck in a loop and then transferred. Or, PolyAI mistakenly admitted a non-authorized visitor.



Testing Tips

There are two types of testing. We strongly recommend keeping these testing types distinct.

Functional

Does the bot behave as expected?

Use the PolyAI UAT sheet to validate all test cases and annotate them using the included scoring scheme on the next slide.

When PolyAI update the agent, we will let you know which cases to re-test.

All cases should be marked as 'As Designed' for Envera to approve go live, unless we mutually agree.



Edge Case

Can I break the bot?

Run some edge case testing, to ensure the agent behaves appropriately.

We don't look for perfect pass rates here, as they should be rare cases, but want to offer Envera comfort in deploying the bot.

Deployment Process: UAT ②



During UAT

Keep the testing sheet open. Check in frequently and ensure the Agent Designer and engineer are live commenting on test cases and listening to calls. You can use PolyAI ticket numbers to reference issues. You can also prioritise them in real time.



There is more advice on how to run the day to day of UAT on page X.

The end of UAT

If you have provided good release notes, active check-ins and explained issues, as well as managed a bug tracker, you will likely have an easy time getting a client to agree UAT is passed, or at least that remaining issues are not go live critical. Use the bug tracker format to prioritise, demonstrating progress.

If you are struggling, try and set a go live 1-2 weeks out. Book daily meetings to rally the client into going live, and plant the suggested date as immovable. If you invite third party vendors, and the people you are relying on to go live, it will create pressure on individuals to hit that go live date.

Deployment Process: UAT ③





Dashboards

During UAT, you should monitor that metrics are being written correctly, and added to Luzmo. You might need to ask the engineers to ensure metrics are added to Luzmo.

Advice on building dashboards can be found on page X.

Deployment Process: Pre-Go Live



Once the client says 'Go'

The engineering team will run an HLD meeting. The PSM is optional, but this is a critical internal governance step. Be sure an engineer has written the HLD document in advance of this meeting.

It is important to then **test the live APIs**. This may require setting up a call with the customer to do an end-to-end production test. You should be sure live APIs work before going live.

Give Access to Dashboards



1. Ask the client for a list of emails who will need access to dashboards
2. Post the list of emails and the region in #studio_portal_access_request channel **and** email the same list to user-access@poly-ai.com. Wait for positive acknowledgement.
3. Add the user to Agent Studio. [🔗](#)

All clients are added to jupiter.polyai.app / jupiter.uk-1.polyai.app, regardless of their project setup. Give them access to custom analytics only on day one. No client should be given access to platform.polyai.app.

Give Access to Statuspage

Instructions are found here [🔗](#)

Deployment Process: Go Live



Go Live Email

You must send an email to the client's main stakeholders, cc Nathan *and* legal, and confirm go live. [🔗](#)

Post in #general in Slack

Celebrate the team with details of the project, a call recording and more in the #general channel in Slack.

The Bridge Call



To go live, you should join a longer (2 hour) call with the client to pull the trigger and monitor live calls. This is important as **you want to be the one to play the first good and the first bad call**. If you own and lead this meeting, championing how good the bot is, and mapping improvements onto the bug tracker, you own the narrative. Think of it as an initial hypercare meeting.

Daily Hypercare meetings

Agree you will run daily meetings for 1 or 2 weeks as soon as possible. The sooner you plant the seed about the duration of this process, the more likely it is to stick. Otherwise, they may extend it forever. In these meetings you should:

- Show and add to a bug tracker
- Send release notes to the clients for every push to live — if you are doing work **show it off**
- Use the dashboard to show KPIs
- Play live calls, including good and bad ones. Pick calls that aren't catastrophic but show out of scope topics or issues you know can be resolved, this helps build trust

Deployment Process: Shifting to BAU



CSM Increasing their Role

As the system is now build, the CSM will increase their role, leading more meetings. They should have been present throughout the build, but now will take an active role in ensuring the customer is fully enabled for BAU, and is confident about their ROI.

Criteria to Exit Hypercare


This criteria is in the kickoff deck, and should be adhered to.

Support

For detailed information on the support process, you should look at the support playbook. [🔗](#)

🔄 For smaller accounts, we are implementing a credit system, reflected in the diagram. Requests for changes will use a credit allowance. This is managed by the CSM and support.

Deployment Process: Shifting to BAU




QBR/MBR

For clients who are larger, there will be a monthly recap meeting set up by the CSM. The CSM and PSM work together to collate the releases from the past month, metrics, and suggested improvements.

Please see CS guidance on how to lead these meetings.

Upsell / Renewal Procedure

Docs coming soon.



In Depth: The Kickoff Meeting

The next slides will run through how to run a project kickoff. This is the PSM's first meeting with the client.


Overview

Introductions

Introductions

You'll be assigned a dedicated support team for the duration of the design and build, to long term support.

Your First Point of Contact



Rob Baron
Product Solutions
Manager

Eshan Singhal
Software
Engineering Lead

Your Development Team

Dialogue Designer
Responsible for the sound and customer experience of your virtual agent

Engineering Team
Runs integrations, telephony, and ML and AI training

Customer Success Manager
Leads steady state support for your virtual agent after go live



The introduction slide is straightforward.

The PSM is the main contact for the deployment, and the client should reach out to you primarily. That avoids miscommunication. It's also a good time to get a single point of contact from the client.

You'll want to introduce anyone else on the call in addition to the concept of a Agent Designer and a customer success manager — there for long term support.

You must change this slide depending on the project. Some will not have a CSM, so you can remove that.

In Depth: The Kickoff Meeting

You should walk through a very high level view of what the agent will do. What is it meant to do from a high level — deflect calls from a contact centre or resolve a call type?

Read the order form before the meeting and be confident in what the customer is looking to do.

There will be a more detailed scope run through, so no need to dwell on this for long.

Overview

Your Virtual Agent

Your Virtual Agent

We build in phases, so need to align on what we are launching together first to ensure maximum impact

When does your pool open for the season?

Frequently asked Questions

Answer customer service calls for simple requests, such as parking, opening hours, and directions.

Can I speak to Housekeeping?

Routing to Departments

Helping your PBX operators route calls to the right team

I need to speak to my host Parry Thomas

Routing to Staff Members

Recognise staff member names, and route the customer directly to them



In Depth: The Kickoff Meeting

The build process is four stages and you'll walk through them one by one.

Starting with design, you can describe this as when we *finalise the customer experience of the agent*. Our designers will consume the FAQ information provided and make the agent come to life.

We'll run through the flows, kickoff telephony integration and the client will get to choose their voice.

This is also the final opportunity to make any changes to what you want to build before we hand over to engineering.

Build Process

Build Process

We follow a tried and tested, four stage process that results in a highly successful voice assistant extremely quickly.

Design

Build

Test

Go Live

Build Process

Design

Build Process

Design

Build

Test

Go Live

PolyAI Deliverables

1. Provide technical resource to work with you and your vendors on integrations
2. Agree and finalise the scope
3. Dialogue designer builds the call flows, writes the content and if applicable, trains the voice actor

Client Deliverables

1. Structured FAQ data
2. Complete API documentation and testing environments
3. Complete SIP workbook
4. Provide contacts for requested personnel including a project manager and introductions to key vendors, including a telephony vendor
5. Select a voice
6. Confirm and lock down the virtual agent scope — this may be the last chance to add features until after go live



In Depth: The Kickoff Meeting



We often describe the ‘build’ stage as a quieter stage where the PolyAI team are ‘heads down’.

We’ll still want the client to support on weekly meetings and to motivate their telephony supplier.


It’s good to explain how we’ll still fill this period by introducing the client to:

- how to interpret dashboards
- how to talk about our agents with contact centre staff
- talk about design methodology
- introduce them to self serve tools
- introduce the BAU support process.

 All of these are coming in new decks in collaboration with CS soon, but even though we don’t yet have content, you should cover off these topics in this period.



In Depth: The Kickoff Meeting



We like to describe UAT as *where the fun starts*. This is what most of our clients enjoy, getting to experience the PolyAI virtual agents for the first time. For this section of the deck, read through the deliverables and explain we'll be making updates *nearly daily* and will be really high touch.

You can describe how we'll provide *a comprehensive test plan to ensure the agent launches well*, as well as a *guide for contact centre staff*, which is the UAT kickoff guide described later.

In Depth: The Kickoff Meeting

It's good to give a brief overview of what to expect on go live day.

This slide includes four examples of how we might roll out. Ultimately, we want more calls as quickly as possible, so build confidence with the client so that they agree ideally to numbers 1 or 2 on this list.

Familiarise yourself with projects that launched in these ways.

Build Process

Go Live

Build Process

DesignBuildTestGo Live

Staggered Rollout

The system can rollout with a percentage of incoming calls. Different clients launch in different ways, for example:

1. **100% on Day One** - some of our customers see the benefit of learning from a high volume of calls
2. **Test on Day One, Launch on Day Two** - for those with new integrations or telephony integrations, it might be safer to start with a small volume on day one, before moving to 100% on day two.
3. **Slow ramp up** - Other customers increase the % every day. We usually aim to be at 100% at the end of week one
4. **Launch per business line** - Finally, it may be possible to launch with a subset of customers by enabling only a single inbound phone number or subset of customers, like existing customers but not new ones.



In Depth: The Kickoff Meeting

It’s important to highlight that we are still closely involved post launch. For hypercare, a usually two week period, we’ll have daily catchup calls. We’ll journey map by playing real calls.

Nathan likes to make a joke here: *we’ll play you great calls and also the really really terrible ones, because there will be one upset person, and we should learn from it.*

This helps set expectations.

Hypercare is not infinite, describe it as **two weeks**. There are criteria listed on the slide for how to pass hypercare. During hypercare, which is described on slide X, you should follow this checklist.

Try not to describe this as the ‘end’, it’s when we decide to move on to the next phase.

Build Process

Go Live

Hypercare

Virtual agents aren’t perfect on launch, because real customers will always behave differently than test cases. For around two weeks, we’ll work closely in a period of ‘hypercare’, to iterate on the agent. There will be:

Daily Catch-up Calls

At the start of a rollout, we suggest daily catchup calls to review live calls, metrics, dashboards and any seen issues

Journey Mapping

Often it is useful to review unsuccessful calls, to scope out new use cases that may not be supported and understand unexpected customer inputs

Support Introduction

PolyAI to agree and share it’s SLAs for feature requests, in addition to support numbers, ticketing, and communication

Hypercare is considered complete when we meet the following conditions:

-  We are live with 100% of all addressable callers
-  We are within 10% of our pre-agreed success criteria
-  PolyAI has provided a support introduction and contact
-  You have access to all dashboards and available self serve tools

Once complete, we’ll move to begin work on the next phase of your virtual agent build. We’ll discuss this more closer to go live.




In Depth: The Kickoff Meeting

Make sure the client knows that we are going to project manage for them.

We'll provide a detailed gantt chart which is a live website they can follow.

This will come later in the process, towards the end of the design phase.



In Depth: The Kickoff Meeting

Build Process

Milestones

Milestone Dates

We have planned with the following milestones in mind. If you're happy, we'll build our full timeline around these dates.

22nd November	●	Build Stages Complete
29th November	●	Testing Sign Off
4th December	●	Go Live
4 weeks after go live	●	Hypercare Period

In the mean time, to comfort the client, provide some generic milestone dates.

This is your chance again to reset expectations. Set the go live for something probably a tiny bit longer than they expect, and that way you can move it forward on request.

In Depth: The Kickoff Meeting

There are two slides to talk about reporting.

Reporting and Metrics

Dashboards and API

Reporting and Metrics

📊

Live Dashboards and Visual Analytics

1. KPIs accessible to all stakeholders
2. Track customer behaviour
3. Explore the next automation opportunities
4. Export as Excel, CSV or PDF

🔌

Conversations API

1. Lists metrics for every conversation
2. Can be ingested into your own data lake
3. Includes in progress calls

The first shows *how* we will report data. There will be dashboards, usually updating every few hours which will be *where you'll spend most of your time monitoring performance*. You can export as an Excel, CSV, or however you like.

Make sure they know we also *have an API so you can import the data into your existing tools like PowerBI or Tableau*. This usually gets a positive reaction.

Reporting and Metrics

Dashboards and API

Measuring Success

Overall Objectives

Reduce calls reaching in-house staff, to free them up for the tasks in front of them.

Eliminate missed calls and bookings.

Improve customer service by answering the phone immediately.

Technical and Contractual Success Measures

Eligible Call Conversion | 43% Target

The percentage of eligible booking calls that result in a booking

Eligible Call Conversion = Calls where a booking is successful / Calls where a customer starts booking, there is availability returned from Opentable that the customer accepts, and PolyAI is authorized to book as per the business rules agreed (not above maximum party size)

The second slide is either the contractual or proposed success criteria.

Get alignment on this — this is what the contract says, this is the definition. Does everyone understand it?

You'll later add this to the dashboards.

V2.0.1

3 DEC 2025

37

In Depth: Scope Run Through

The scope run through is usually the second meeting you'll have with the client. You'll talk through either how to provide PolyAI FAQ information, or go step by step through the flows. This is difficult, technical, and involved.

Coming Soon

In Depth: Implementation Review Board ①

In order to maintain project quality, before going into UAT, there will be an 'implementation review board' meeting. In this meeting, held at least 7 days before the planned client UAT kickoff, a set of requirements will be presented to the most senior engineer, Agent Designer and PSM in the team. As of Feb 2025, these are Eddy, Arkadiusz and Nathan.

The meeting agenda is as follows:

1. Brief overview of project functionality
2. Walk through completed QA sheet
3. Show metrics configuration
4. Run through checklists
5. Approval by Eddy, Nathan, and Arkadiusz

The outcome of the meeting will be one of three things:

Approved - The project can proceed to UAT on time

Approved with Conditions - The project can proceed to UAT on time provided that a set of conditions are met asynchronously, before the UAT Starts

Rejected - The project cannot proceed to UAT and the Implementation Review Board must be run again.

In Depth: Implementation Review Board ②

The checklists that require completion are in Notion. Use the template to create a new Notion page to prepare for the meeting.

The checklists are split into three categories.

QA

Containing evidence that QA was completed to a high standard and all aspects of the virtual agent were tested. Includes some specific frequent points of failure. Also includes evidence of metrics being fully and properly implemented to given standards

Project Behaviour Standards

A list of behaviours that are expected to be present in all virtual agents. All standard project behaviours should be tested as part of QA

Logistical / Process

A list of process requirements that should be complete by the time the board is met, including creating the UAT contents and ensuring metrics are in Luzmo.

Table

QA

Aa Item	Completed	Evidence
100% of functional test cases pass	☐	
QA includes metrics, and all test cases pass	☐	
QA includes out of hours behaviour, and all test cases pass	☐	
QA includes API Errors, and all test cases pass	☐	
QA includes cases for disambiguation notes, which must understand both options provided word for word, in addition to yes/no/don't know	☐	

+ New page

+ Add a property

Comments

Add a comment...

Meeting Date: X

Meeting Scheduled

To-do

Meeting Scheduled

Rejected

Complete

Approved

Conditionally Approved

In Depth: The UAT Client Kickoff

This is how you introduce UAT. You'll be setting expectations here, so explaining this well is critical.

The first slide of the UAT kickoff deck contains the test phone number.

It also sets expectations on a few common misconceptions about UAT. It won't be perfect. The voice may change or be generated for the first time causing latency. And handoffs are usually disabled in UAT.

On this slide, you should also add if the agent is using sandbox APIs.

Envera Test Brief

Summary

Summary

PolyAI have created a virtual assistant designed to screen community visitors and either grant access or transfer to human agents.

Things to Note

1. The data used to train the system will rapidly improve after your testing and real customer calls. There will be a short ramp up in accuracy.
2. There may be some phrases with an unusual or robotic voice. Please alert PolyAI if you discover these.
3. This test number will not transfer to human agents

Test Phone Number
0800 00 1066



In Depth: The UAT Client Kickoff

Pradco Test Brief

Summary

Testing Tips

There are two types of testing. We strongly recommend keeping these testing types distinct.

Functional

Does the bot behave as expected?

Use the PolyAI UAT sheet to validate all test cases and annotate them using the included scoring scheme on the next slide.

When PolyAI update the agent, we will let you know which cases to re-test.

All cases should be marked as 'As Designed' for Envera to approve go live, unless we mutually agree.

Edge Case

Can I break the bot?

Run some edge case testing, to ensure the agent behaves appropriately.

We don't look for perfect pass rates here, as they should be rare cases, but want to offer Envera comfort in deploying the bot.

It is extremely common for clients to misunderstand UAT and cause us significant problems.

Therefore it is very important that you **familiarise yourself with and understand the different types of testing**, and the client should do.

We want to be very distinct and serious about the difference between functional testing, which will use our prescriptive UAT sheet, and edge case testing, which is when a client tests in a less structured way.

You should explain that **functional tests are the test cases we give you that are scripted and should work every time**. After they are complete, then the team can do edge case testing, which is their own exploration.

In Depth: The UAT Client Kickoff

We use a scoring system for functional testing.

After this deck is complete, you will walk the client through the testing sheet. See pages 20-22 for more details.

You should also talk through the scoring system and make sure they understand it.

If the client understands this, and accepts we can go live provided all test cases are in the top three scores, you will have a much easier time convincing them to launch.

Scoring Functional Testing

This is how to score a functional test you complete.

As Designed - Perfect

The outcome of the call was as designed. For example, a visitor who should be granted access was granted access. The user was understood first time, and the CX was good.

As Designed - Imperfect

The outcome of the call was as designed. For example, a visitor who should be granted access was granted access. However, the experience was not perfect. The outcome was correct, but the user may have had to repeat themselves.

Fail - Good UX

The outcome of the call was not as intended. For example, the agent transferred a user to a human operator despite the fact the visitor was authorized to enter and the API provided correct information. However, the user experienced a seamless experience, such as a fast transfer to a human.

Fail - Bad UX or Incorrect Admission

The bot did not behave as intended, and the experience was bad, for example, the user was stuck in a loop and then transferred. Or, PolyAI mistakenly admitted a non-authorized visitor.



Ask how to reset
want to do it on t
Ask how to reset
want to do it on t
Ask how to reset
find your verific

In Depth: The UAT Client Kickoff

Finally, use the final slide in the template to create some conversation scripts to explain how to test. You can show that you can slightly go off topic, or avoid loops.

You might also like to play a few live calls or **make a live test call together**.

Following this, walk through the actual UAT sheet. Use the example on page 21.

Envera-PolyAI UAT Phase 1 External						
File Edit View Insert Format Data Tools Extensions Help						
Q Menus 100% 103 Default... 11						
A1 Premise number: +1253878113						
1	Prerelease number: +1253878113	Community ID: 3460801401				
2						
3	Test Case	Description	Correct Visitor Context / Info	Instructions	Expected Outcome	Test Score
3.1	Visitor Drive-Off - Pre DL Scan	The visitor drives away before completing the PolyAI flow	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. When asked to insert DL, remain silent and do not insert DL x2 3. When asked other questions, remain silent x2	1. PolyAI ends connection 2. PolyAI calls agentresult with validated = false	As designed - imperfect
3.2	Visitor Drive-Off - Post DL Scan	The visitor drives away before completing the PolyAI flow	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Insert DL (use surname only) 3. Remain silent x2	1. PolyAI calls agentresult with validated = false	As designed - perfect
3.5.1	Early Scan Visitor Name Match With Confident Resident Name Match - Positive Match	Match Visitor from DL, Match Resident Correctly	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Insert DL (use surname only) 3. Respond to Company Question as Individual 4. Give Visitor Full Name correctly 5. Provide Correct Resident Address 6. Provide Correct Resident Full Name	1. PolyAI calls agentresult with validated = true and correct details (You are let in)	As designed - imperfect
3.5.2 (and 6.12)	Early Scan Visitor Name Match With Confident Resident Name Match - Positive Match no DL	Match Visitor without DL, Match Resident Correctly	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Either no DL, Scanner or remain silent and do not enter DL 3. Respond to Company Question as Individual 4. Give Visitor Full Name correctly 5. Provide Correct Resident Address 6. Provide Correct Resident Full Name	1. PolyAI calls agentresult with validated = true and correct details (You are let in)	As designed - imperfect
3.5.3	Early Scan Visitor Name Match With Confident Resident Name Match - Negative Test Resident Name	Match Visitor from DL, Give Non-matching Resident Name	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Insert DL (use surname only) 3. Respond to Company Question as Individual 4. Give Visitor Full Name Very Incorrectly 5. Provide Incorrect Resident Address 6. Provide Very Incorrect Resident Full Name until transferred	1. Call is transferred 2. PolyAI calls agentresult with validated = false and collected details	As designed - imperfect
3.5.4	Early Scan Visitor Name Match With Confident Resident Name Match - Negative Test no DL, Visitor Name	Give Non-matching Visitor and Resident Names	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Either no DL, Scanner or remain silent and do not enter DL 3. Respond to Company Question as Individual 4. Give Visitor Full Name Very Incorrectly 5. Provide Incorrect Resident Address 6. Provide Incorrect Resident Full Name	1. Call is transferred 2. PolyAI calls agentresult with validated = false and collected details	Fail - Bad DL or visitor mistakenly admitted
3.5.5	Early Scan Visitor Name Match With Confident Resident Name Match - Negative Test Fuzzy Resident Name	Match Visitor from DL, Match Resident Correctly with CLOSE alternative Resident Name	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Insert DL (use surname only) 3. Respond to Company Question as Individual 4. Provide Correct Resident Address 5. Provide Resident Full Name as Sheldon Cooper	1. PolyAI calls agentresult with validated = true and correct details (You are let in)	Fail - Bad DL or visitor mistakenly admitted
3.6.1	Early Scan Visitor Name Match with Confident Resident Nickname Match	Match on nickname for the resident (Raj vs Rajesh)	Visitor - Amy Farrah Fowler Resident - Rajesh Koothrappali Resident Address - 321 Rocket Science Lane	1. Initiate Event 2. Insert DL (use surname only) 3. Respond to Company Question as Individual 4. Provide Correct Resident Address 5. Provide Correct Resident Full Name as Raj Koothrappali	1. PolyAI calls agentresult with validated = true and correct details (You are let in)	As designed - imperfect
3.6.2	Early Scan Visitor Name Match with Confident Resident Nickname Match - Negative Test	Do not match on the	Visitor - Amy Farrah Fowler Resident - Sheldon Cooper Resident Address - 333 Rocket Science Lane	1. Initiate Event 2. Insert DL (use surname only) 3. Respond to Company Question as Individual 4. Provide Incorrect Resident Address	1. Call is transferred 2. PolyAI calls agentresult with validated = false and collected details	As designed - imperfect

Getting The Most Out the System

In addition to functional testing, we recommend some creative testing, to talk like a real customer would, slightly changing topic midway.



Hi, the colour of my camera is all off?



If your night photos come out all black, or if your day photos are pink or white, it might be time to reset your infrared filter. Can I tell you how to do this?



They're pink!



Yes, if your daytime photos are appearing pink, it indicates an issue with the infrared filter. Would you like me to explain how to reset it?



In Depth: UAT Day to Day

Coming Soon

In Depth: Release Log

When we push to live, we must update a release log. This is a compliance requirement, adds professionalisation, and is good defence against a client. This must always be updated. **It is owned by the PSM for build and hypercare**, and then can be transitioned to CS or the scrum master.

The release log template  table includes a release date, a deployment owner, a version number and a list of changes. They are stored in a shared Google Drive. These should:

- Be in readable english and generally non-technical
- Include ticket numbers, including the support desk external ticket number where relevant
- Have evidence of QA

There is a separate box for QA, which is our evidence to a client that we have done due diligence and tested the virtual agent. It is acceptable to put minor known issues in here.

An internal QA doc can be included that the PSM can later refer to as evidence.

Regardless of the size of the push, a release log must be added. There is a Google Docs shortcut, just type *@Release Log* to auto-insert the table.

To share this with a client, either share the live Google Doc or export as a PDF and send periodically.



Example Client | Virtual Agent Release Notes

Document Owner:
Project Name:

February 12th 2025

Release Date	Deployment Owner	Deployment Hash / Version Number
February 12th 2025	Archawin Wongkittiruk	Previous: fcb046254919f10136d1f078132c06314a71d742 New: 252fa9dafdff0c3f9aefd8c73901394e18a617c

Changes

Change Description	Ticket Number (include support desk ID if applicable)	QA Pass?
Reverted Dogs/Gates safety issue change. No longer force transfers for callers who refuse to answer if dogs are kept behind locked gates	Reverting UTIL-1741	Y
Add metric for ineligible account transfers	UTIL-1731	Not Tested
Resolve a case where calls could crash in a specific scenario of a spanish speaker calling out of hours, and triggering the payments intent	UTIL-1773	Y
Ask 'is there anything else I can help with' after informing a user the contact centre is OOH	UTIL-1522	Y

Quality Assurance / Testing:

QA Owner: Oliver Shoulson
Known Issues/Fails: New metric will be validated in live, check by 2/13 EOD
Notes: None
Link to External QA Doc:  Authentication-InternalQA

Day to Day Team Meetings and Activities


Coming Soon

In Depth: Dashboards | KPIs Dash

Good analytics dashboards are vital to your projects being a success. There is no compromise in the detail and professionalism. This section will highlight the minimum things to look out for when creating a dashboard, based on the Fogo de Chao dashboard.

Neat, symmetrical, layout

Color palette matches the client's brand standard




The most important success rates for the client are at the top

Trends are useful, weekly for overall success and daily for 'at the moment' check-ins. The top graph could also be hourly.

In Depth: Dashboards | KPIs Dash

Easy to understand, grouped categories of call reasons. These aren't raw metrics, but understandable, natural english categories



Containment is measured per-topic

Most important to every client is knowing how to get better. This pie chart shows it

In Depth: Dashboards | Raw Call Log

The raw call log is a vital dashboard to see detail, per call information. It's a simple table to make.

Raw Call Log

Download report

Last 7 days (includes today)

Fogo de Chao - Lo...

Topic

True

Successful Booking

Booking Outcome

Restaurant Name	conversation_id	Date and Time	Questions Asked	Reason for Transfer	Handoff To	Booking Initiated	Successful Booking
Fogo de Chao - Los Angeles	KA89865f3a-deab-4611-bee9-eac19c1a1348	07-07-2025, 02:32:41	reservations-make_booking	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KA81f1f031-d063-45d3-a5f8-1e94e0c94a9b	01-07-2025, 02:56:57	speak_to_human	restaurant_max_group_size	SALES_MANAGER	True	N/A
Fogo de Chao - Los Angeles	KAc d8e0817-e57f-46de-a737-fc17e6c2ea65	02-07-2025, 03:39:05	reservations-make_booking	restaurant_max_group_size	SALES_MANAGER	True	N/A
Fogo de Chao - Los Angeles	KA5a713281-5580-4a7a-8dc2-9389f82fc396	02-07-2025, 23:24:18	reservations-make_booking	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KAf0e8339e-b892-4c34-afe7-8aa41ea09749	04-07-2025, 23:42:19	reservations-make_booking	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KAe47acf1c-fc9c-43b7-a7a9-4e59f52e9a99	05-07-2025, 01:41:02	reservations-make_booking	restaurant_max_group_size	SALES_MANAGER	True	N/A
Fogo de Chao - Los Angeles	KAa712ccb7-734f-492b-9020-bb71b5434232	03-07-2025, 20:53:13	speak_to_human	speak_to_human	STANDARD	True	N/A
Fogo de Chao - Los Angeles	KAadd59d7b-45ae-4cfd-8b33-1b03b5f1719f	05-07-2025, 04:07:05	site_opening_hours	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KA0ca1900f-94ed-452b-b0b1-a6ca388dd56c	03-07-2025, 23:01:43	reservations-make_booking	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KAb56a0f4d-64f3-4f9a-b25f-b4724c3e0786	05-07-2025, 04:33:05	reservations-make_booking	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KAef21a6d8-6de2-4817-ba18-3a96a97c671d	04-07-2025, 02:24:39	reservations-make_booking	N/A	N/A	True	N/A
Fogo de Chao - Los Angeles	KAfcee1845-d4e7-4fdf-a8a3-3d00101f7e79	04-07-2025, 03:58:52	reservations-make_booking	make_booking_in_the_past	STANDARD	True	N/A
Fogo de Chao - Los Angeles	KA04044cd0-688e-4c89-8d70-ade1a6555bd3	06-07-2025, 23:45:33	speak_to_human	speak_to_human	STANDARD	True	N/A


Filters to highlight specific calls for deeper analysis

Virtually every important metric is in this table (hidden)

Dashboard set to link directly to conversations review

In Depth: Dashboards | Industry Specific

Many clients will have very specific requirements or industry specific data. **Create a different dashboard for this.**



Restaurant excess demand dashboards show where PolyAI is turning customers away because there is no availability

Going on Holiday

When you go on PTO, it is vital to create a good handover document so that your cover knows exactly what to deliver.

When taking PTO:

- Create a Notion page
- List all your projects
- List deliverables that **must** be complete at the end of your holiday, in detail, including the tools that another PSM will need to be able to complete that task
- List the main contacts for each project
- Invite your cover to scheduled meetings
- Send an email to the client to introduce them to your cover

Moxie Pest Control

Status: **LIVE with Stafford (Moxie Pest Control) & Kansas City (Joshua Pest Control); in hypercare; Targeting go-live at 9 more locations late next week**

DD: @Sophie Moradi

Engineer: @Archawin Wongkittiruk

Coverage: @Becky Kershaw @Rob Baron

POCs at Moxie:

- Steve Woeltge (their VP of Customer Care and our main Champion)
- Treasure Ford (Project Manager, main poc)
- Alyssa Mancini (scheduling expert, also knows a LOT about their product)
- Ryan Turner (their IT guy. Gives us the API keys we need)
- Chris Pettitt (handles routing/telephony on their end)

At a high-level:


- 15 minute daily hypercare check-in
- Check hypercare sheet for issues reported by Moxie and post in channel if the team isn't aware and getting an update on work in progress to update the Moxie team that issues being raised are being resolved
- Hypercare bug tracker:
https://docs.google.com/spreadsheets/d/134ylsl0Wlx6gUfFJaz_TDIV7f0evpdUOLIQSdt7wcsU/edit?gid=1193964895#gid=1193964895
- Ensure live call reviews are happening
- Rollout readiness assessment/push to live

Work that needs to be done next week:

- "Phone Number Not Recognized". This issue is happening where we do get a phone number, accurately validate it, however the response we get from the API indicates that there is no account match received. Upon investigation these cases that are being noticed at Stafford



Never book holiday over a UAT or go live period. As you have a timeline in smart sheet and a go live date range, you know when not to book holiday.



PART TWO

Managing Clients

Dealing with Objections

This section will run through how to manage common questions and objections that happen during the build process and beyond.

Something a client might say

A good way to approach the answer

Some additional notes to help you manage an objection

Early Project Objections: Scope

Can I review the design of a virtual agent before you build it?

Yes, we'll have a scope meeting to run through the flow logic. We'll also be able to share the knowledge base the virtual agent is using to answer the caller. Please share your brand guidelines to help us, but we'll take the first shot as our team must get the foundation working first.

How many FAQs can I have?

There is no hard limit on the amount of topics. And each topic will be able to answer many different questions, depending on how a caller asks it. However, to start we recommend no more than 50, to prevent overlap. We'll then track all of the out of scope customer requests and decide which more to add.

Why should I pay / take on risk to send SMS?

Directing a customer to another channel through SMS is a very effective way of gracefully deflecting customers. In one project with over 30% total containment, **calls where an SMS is sent are contained over 85% of the time**. This goes for integrations too — when PolyAI *does something* for a customer, they have a better experience.

Early Project Objections: Scope

Can I start with a small number of FAQs and add more later?

Yes, but it's important we know a broad range of things a caller might be asking about so we can deal with it properly. Do you have data on how many times questions are asked, so we can prioritise by volume?

Early Project Objections: Voice

Can I use my own voice actor?

We have a large selection of voices to choose from. Some of these are our own PolyVoices, which you won't find anywhere else and are trained on our own data and voice actors. Others are 'off the shelf' voices. To train new voices, we pay significant license fees, and conduct studio recording. We aren't quite in the position to clone any voice at the moment. If you are unsure, pick a PolyVoice.

Early Project Objections: Metrics

Can I measure customer satisfaction / CSAT?

Yes, there are a few options.

1. We can do a voice survey, but on contained calls only, because we don't want to put a survey before the customer is transferred to a human, that's annoying. These results will appear in PolyAI dashboards.
2. You can provide us a survey link and we can send it via SMS on the end of every call. These results will appear on your end.

Generally, we advise CSAT questions ask specifically about the 'virtual agent' as in our experience customers tend to use surveys to complain about *anything*, not just the virtual agent.

If setting this up is going to take time, we can use other proxies for CSAT like call length and containment for the specific flows.

Early Project Objections: Content

I want the agent to introduce itself as a virtual assistant.

We tend to suggest not doing this. People generally don't like voice assistants, and we need to have one chance to build some trust with the caller without them knowing. If a caller asks, of course we'll tell them we are a virtual agent but we see about a 25% uplift in containment when we don't declare we are AI. We don't want to fool anyone into thinking they are talking to a human, we want them to forget that they aren't.

I'm worried about customer experience so I want to change this FAQ so it transfers more easily

Yes, we can do this. We will warn you that we find when you offer a transfer, the customer nearly always says yes. Of course, particularly unhappy customers or vulnerable ones will always get transferred.

Note: use discretion. Are we meant to contain this call topic or is it completely out of scope? Usually common sense will tell you when to say yes or no to this.

Early Project Objections: Content

Can I use the bot to market my brand? / I want to use this slogan in this FAQ response

We're happy to help with that, we think that engaging customers is most important so don't want to put any off-putting marketing, or anything that sounds like it a robot. We want customers to pay attention and have a natural conversation, so usually avoid it. How do you tell your contact centre staff to promote things, do they say the same? If not, the bot shouldn't.

How are you going to deal with vulnerable customers, I want to transfer them.

We can detect vulnerable callers through keywords and the topic of the call to transfer them quickly. Can you give us what topics you want immediately transferred? We can test them in UAT

Early Project Objections: Call Listening

Can I listen to calls?

Yes. After go live, we'll have daily catch up calls and we will play you both the best and the worst calls so we can do journey mapping. Then we'll give you an API and platform so you can pull the audio yourself, and also mark calls for re-training.

Note: We should be careful here. We do allow customers to listen to calls, however the perception of PolyAI can change. Especially when we move on to other projects, customers may feel their agent is degrading. **Make sure when you introduce call listening, you do on a screen share, and have a session to consult them on how to listen productively and find both good and bad calls.** Getting this wrong will make your work harder.

UAT / Live Call Objections: Quality of Recognition

It won't understand me / every time I say X it doesn't work

We'll look at all the calls, let's look together. I'll add it to the issue tracker. This is why it wasn't understood...

The best way to handle objections during UAT is to fix it. The client will come across things that aren't understood, but you know they can be fixed fast. It's also good to explain why: look, we were missing this word from the training data, for example.

It doesn't work with background noise

We'll adjust the latency threshold. Sometimes it's a fine balance as if we make it less sensitive to background noise, the agent might interrupt people. We've got everything set to the default, we'll adjust from there when we go live/

UAT / Live Call Objections: Voice

The latency is bad, there are long pauses

We'll make some adjustments. There is a trade off here, the shorter we make the latency the more likely it'll interrupt. I wouldn't worry too much at this stage, we'll tweak and over time you'll see it improving.

Note: The engineering team have a few steps to take here. You should not accept anything more than 3 seconds at an absolute maximum. You should make sure the engineering are not overriding VAD, that they have used the latency graph generator to see if there are low hanging fruit, and also that they have delay responses such as 'one second' and 'umm' and 'i'm just looking that up' enabled. Don't take no for an answer!

Why can't I interrupt the virtual assistant or barge in?

We have this feature, we've learned that agents perform much better when it's turned off. This is because sometimes a customer will interrupt and miss an important part of the agent's speech, or a question the virtual agent needs to ask. We're working on improving it, and will let you know.

UAT / Live Call Objections: Metrics

Does containment include people hanging up? How do I know they are good calls?

Yes, containment is our true, unadulterated future that shows simply how many people are not transferred to an agent. Let's create a success measure ourselves that removes callers who don't interact, short calls, and more.

Note: Often, the contract will contain a defined success criteria. You should make sure dashboards show that criteria beside containment. However, be careful about creating criteria we can't control. For example, a booking success rate should not include cases where there is no availability or we are not permitted to make the booking in the denominator.

The dashboard doesn't match our figures?

The dashboard is sometimes delayed a small amount, I'll look into that. What time range and time zone are you setting on the dashboard?

Note: You may need to reset Luzmo warp. Next, check the time zone on the dashboard. Ask the client to give you their data to compare.

Internally Escalating

There are many slack channels for escalating questions

Sometimes you or an engineer might meet the end of the road in solving a client problem. Look to the right team to help.

#psms-internal, to chat to just the PSMs

#ask-agent-studio, for questions or problems with the Agent Studio web app

#ask-voice, for things related to the voice team, such as latency, speech recognition or speech synthesis

#ask_legal, for questions about the contracts or permissions for your clients

#ask-portal, is for questions about the legacy PolyAI portal for intent based agents

#studio_or_portal_access_request, when you need to give a client access to Agent Studio

#ask-analytics, for dashboard related queries

Escalating Well

You will have difficult projects, and it is important to know how to use your 1-1s, account reviews, and other meetings to escalate and ask for help at the right time.

Which Category Does the Problem Fall into?

The problem does not pose an immediate risk to the task/project/org. You are in control to resolve it.


The problem is site sensitive/serious and there are risks/consequences involved. You may be in control and need managerial support.

Actions to Take


- Inform only
 - Communicate your solution via slack or in a 1-1
 - No actions needed from the manager
 - Let your manager decide whether they want to step in and help
- Escalate quickly and specify action needed from manager
 - Communicate via slack immediately, and discuss further in 1-1s
 - Communicate the solution

PART THREE

Agent Design



Coming Soon



PART FOUR

 Technical Guides

Telephony In General

Telephony will either be connected via SIP or PSTN. Read more on telephony here [🔗](#)

SIP (Session Initiation Protocol)

- PolyAI preferred telephony set up
- This is telephony 'over the internet', connected by a SIP Trunk. A SIP trunk is created by the telephony provider on your project to allow for calls to be routed to PolyAI and from PolyAI back to the client
- Allows for sending of information on call termination and handoff e.g. handoff reason and FAQ asked
- Set up time can vary depending on vendor see the next slides for more information

PSTN (Public Switched Telephone Network)

- Traditional legacy phone network
- Phone to phone with no need or requirement of the internet
- Very fast to set up (1-2 days) but has low capability e.g. cannot send any additional conversation information to a human agent on handoff
- To complete the integration:
 - PolyAI will set up a phone number for the agent
 - Vendors will route their private number to the new PolyAI number
 - Vendors may need to set up a 'secondary' number that PolyAI sends calls back to on transfer



The following guides mention the SIP Interoperability Doc. This is a document that acts as a project plan, run by an engineer. It's worth checking in frequently on the document to make sure it's up to date. You can also show it to key project stakeholders to describe the process.

How to: Generic SIP

Send our SIP Workbook



Interoperability doc created
by Engineering



Telephony Kick Off &
Timeline Agreed

Things to Remember

- Ensure an SoW is in place between the client and the vendor
- Make use of the telephony slack channel and telephony engineers
- Kickoff early
- Ensure dates in the interoperability doc are in the main project timeline
- Have at least one-two weeks between the telephony vendor completing their work and go live

How to: Five9

Email Five9 Team for initial setup



Interoperability doc created by Engineering



Telephony Kick Off & Timeline Agreed

Things to Remember

- Five9 are slow, and an SoW may be required between the client and Five9. You will want to start this progress immediately at kickoff, without delay.
- Five9 will assign a resource, such as a project manager. Keep pushing for a kickoff meeting with them.
- Five9 has a multi-tenant SIP trunk that connects to PolyAI in the USA. Some clients assume that an entirely new trunk needs to be provisioned. **It does not.**
- Five9 will typically send us a SIP document for us to complete.

 Detailed Architecture Information

How to: NICE inContact / CXOne

Email for setup, requesting pseudo numbers



Interoperability doc created by Engineering



Telephony Kick Off & Timeline Agreed

Things to Remember

- There are two ways to connect to NICE. Each client has their own set of pseudo numbers assigned to them by NICE. We can either use these numbers or our *connector auth token* to route calls to a project. Ask your telephony engineer which they prefer, and propose it.
- NICE uses the same IPs across all clients, so we can use the shared nice connector service integration.

 Detailed Architecture Information

How to: RingCentral

Coming Soon

How to: Payments and PCI

PolyAI can take card payments over the phone, using third party PCI compliant systems.

Discussing with a client

We can take payments and make card holds securely using most payment gateways and providers. The payment will be taken in a secure section of the call that PolyAI does not and cannot record. PolyAI will be able to generate a token, or execute a payment. Card payments and ACH are generally supported. Our PCI documentation is available on request.

Payments are generally specified in the order form, along with the payment gateway.

Internally / Technical Detail

We use Twilio Pay to securely take the payment. This service handles the payment so PolyAI is unable to see the private card information. For payment providers that Twilio Pay supports out of the box (e.g. Stripe, Braintree, Base Commerce, Card Connect) we can build immediately. Your engineer will tell you which details we need from the client.

If the payment provider is not supported by Twilio, check to see if it is supported by Shuttle.  If it is not supported by Shuttle, we can pay Shuttle PS for them to build a connector. Shuttle allows us to use additional payment providers with Twilio Pay.

PART FIVE

 Top Tips






be a human

Build a connection with your clients and **humanise yourself**. Make them feel smart, make them feel like they came up with ideas. **Have a story** behind your decisions. **Be vulnerable, you're in this together** with your clients.



avoid conflict

Blame process, not yourself. Be firm when saying no, but try and make it a 'yes but'. Frame your noes as opportunities. **Avoid conflict, especially early.** Usually you can find an agreement, or be vulnerable and make them understand why something might not be so simple.




create a cushion

Pad your timelines upfront, and adjust later to accommodate. Make it clear you are flexible. Emphasise the great services we provide for our clients, our malleability and how custom deployments can be. **UNDER PROMISE, OVER DELIVER.**



make scope creep a win

Don't say 'it's out of scope'. Say, 'let's explore it'. Read back scope status regularly to manage scope creep gently. **Accept a few small changes to avoid the big ones.** Use the scope run through meetings effectively.



perception matters

Make the team seem bigger than it is, especially for large clients. **Bring in engineers and Agent Designers and emphasise them.** Call them leads and talk about their teams. Build credibility with out expertise. On-sites are incredibly valuable, people are much nicer and more reasonable in person. **Travel.**

