Deployments Playbook

SECOND EDITION

A guide to delivering virtual assistants at PolyAI, for the deployments team

V2.0.1
3 DEC 2025

PolyAI

Introduction

PolyAI builds virtual assistants, usually for voice, that empowers callers to engage with a technology that has a bad reputation. At their best, our assistants are indistinguishable from humans and drastically reduces the need for humans in a contact centre, reserving them for the callers that need empathy and tact.

Your job is to deliver these agents, to maximise the quality of them and to be client centric.

Our mission is to make everyday customer service feel superhuman, with voice assistants people actually like.

Your mission as the PSM is to turn this into a reality.

You will make our clients feel excited about the work they are doing. You'll chase ROI for our customers, and even when the measure is blurry, you'll chase engagement for their callers.

This playbook is your guide.

Introduction

This document will use icons and links to take you to further information.

You can click the link icon to find supplementary information, sample documents and code snippets.

This icon signals that this page is something that might soon change, for example, the product is evolving. Check with the team.

PART ONE

Project Delivery

What We Build

Voice Assistants that are Superhuman

PolyAI builds voice assistants. A good voice assistant is indistinguishable from a human for most tasks. The voices are respectful, natural, casual and convincing. They don't sound like robots, adverts, or technology.

We *do* offer text virtual assistants, in chat on the web, via SMS or WhatsApp, or in app. However we build *voice first* — this is because it's extremely hard to convert a chatbot to voice, but easier the other way round. Voice has many nuances such as background noise, crosstalk and bad signal. We therefore do not have any major live chat deployments today.

Voice Assistants that are Revenue Generating and Cost Saving

Our mission is not to replace contact centre agents, but to ensure that callers get through to humans only when they require additional empathy, care and tact. Contact centre agents should take the most interesting and complex calls. They also get the chance to work with us on their virtual agent.

We also can generate revenue. We are careful about how we measure this, as it's often a little out of our control. But we measure how many bookings we make, abandons we reduce, and try to link hold time reductions to revenue generation.

Generative vs Intent Based

Generative

- ✓ A generative AI system answers questions using a knowledge base that customers can edit
- ✓ Conversational flows can be created in PolyAI's Agent Studio, again with an LLM at its foundation
- ✓ Complex flows and high sensitivity ones like 'authentication and verification' flows may be built on PolyAI's prescriptive NLU model and policy framework by PolyAI engineers, but generally not.

Intent Based

- ✓ PolyAI takes a user input and maps it to an 'intent'. The system has prescribed behaviour, so takes a defined action based on that intent, which might be reading out an FAQ, or proceeding through a flow
- ✓ The answers the virtual agent gives are pre-written and fixed

⚠ You should be careful how you talk to clients about this. The exact mechanisms of how we build are not important. **PolyAI uses a hybrid system of both proprietary models and large language models, which are used interchangeably based on the customer's request.** Our systems are predictable and custom.

How We Build

We Build on a Generative AI Platform, called Agent Studio

Agent Studio is our platform and control panel for virtual assistants. It contains the knowledge base, flow design, voice generation, analytics, conversation review and more. Each of these different items can be enabled on a per-client basis.

Our platform has...

A configurable Large Language Model, which may be GPT, hosted on Azure AI, or our in-house LLM, known as Raven. We choose which to use on a per use case basis. Regardless, many client's will understand this as the PolyAI LLM.

A PolyAI Speech Recognition System. Designed and trained in house, this translates audio of human speech into text, which gets sent to the LLM. It's based on a base model from NVIDIA, who have invested in us, and we build on top of it.

A lifelike hybrid Text to Speech model. We offer samples to our clients that might be a mix of a PolyAI trained voice, based on an actor that we work closely with, off the shelf voices, or live recordings. Our client's don't pick, they trust us to use the best hybrid of all three options.

Who Builds

The Deployment Team is a subset of Team Y, a.k.a. The Project Team or the PS Org

It contains Agent Designers, Product Solutions Managers and the Support team.

Product Solutions Managers are the primary client contact for the build of a virtual agent. They own the timeline and the project plans. They build and manage dashboards, analytics, and success rates. They make the client feel good about what they've purchased.

Agent Designers script the agent and prompt the LLM. They are responsible for the quality of the agent and making it as lifelike as possible. They own relationships with the voice actors and produce deep analytics.

The Support team manage our service desk, where live and running clients submit fault requests.

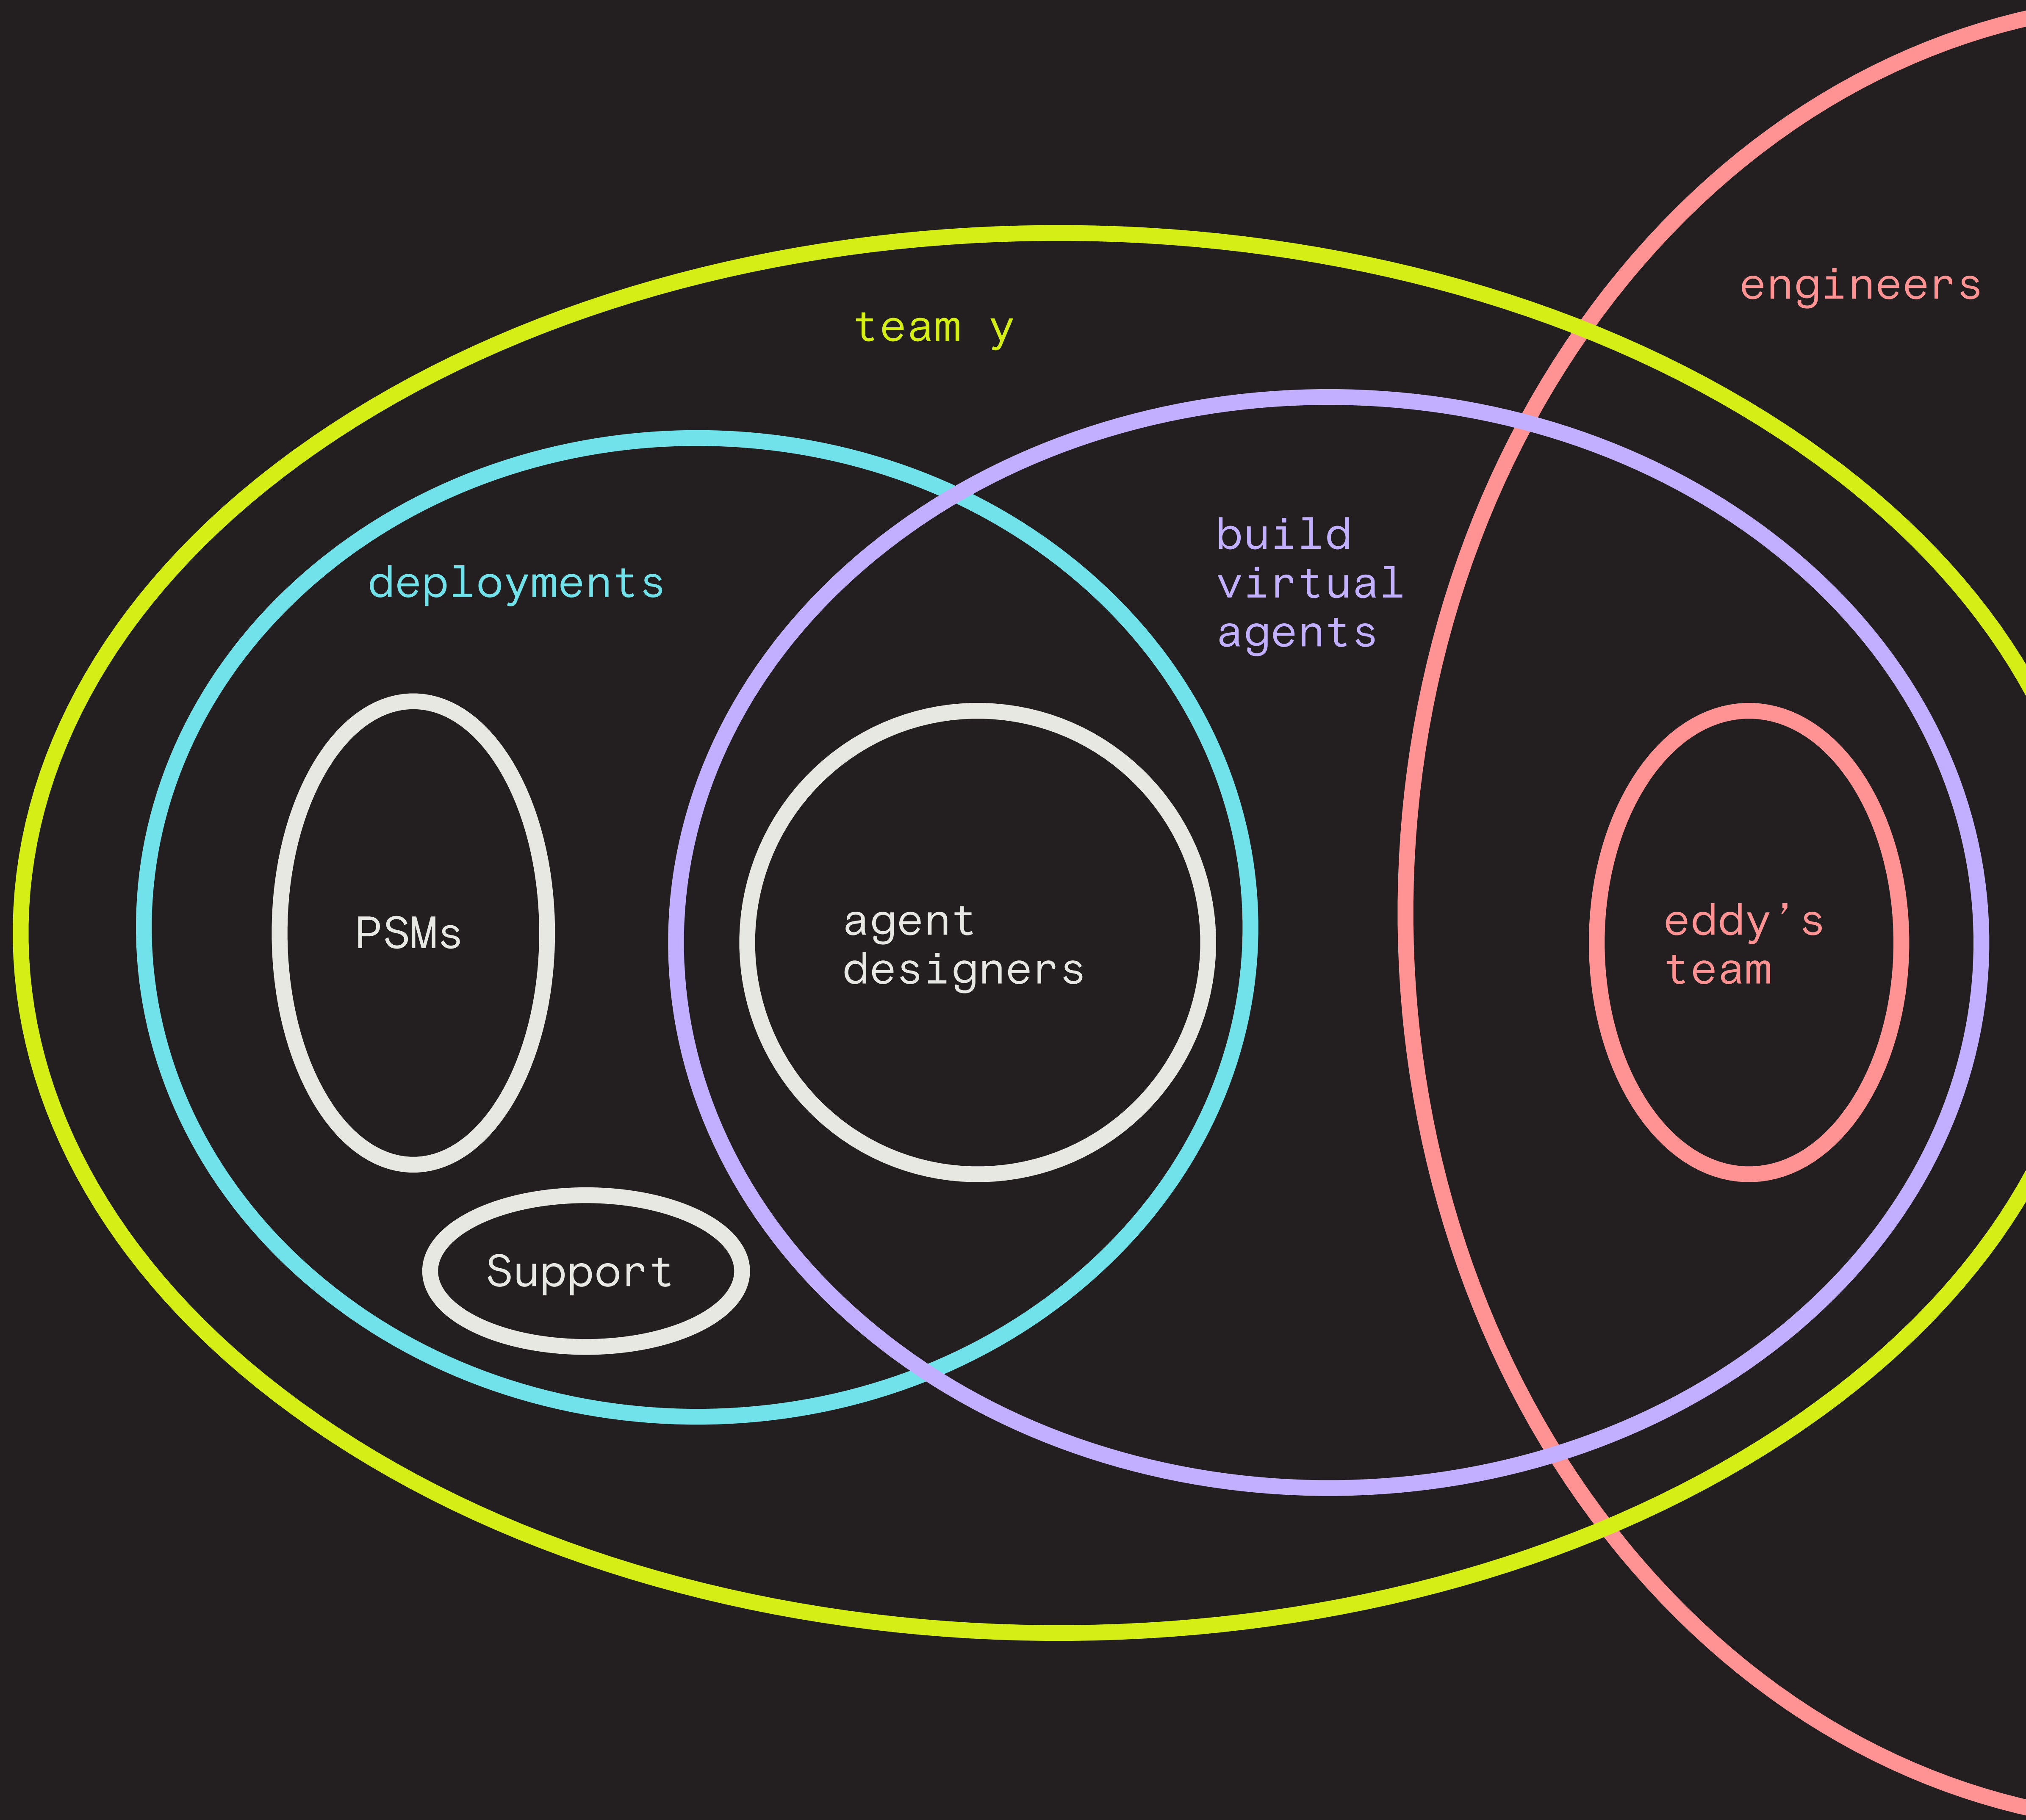

Who Builds

Team Y, the wider Project Team, operates in ‘pods’

Pods contain scrum masters and engineers, who report into Eddy’s org, agent designers, PSMs and a CSM.

New projects and clients are assigned to the pods.

The pod makeups may change, but at time of writing are:

Bao, primarily based in the UK and responsible largely for FedEx

Sourdough, primarily based in the UK and responsible for many European clients

Somun, primarily Serbian engineers but responsible for US projects due to EU/UK data transfer rules

Focaccia, based in the UK with global projects

Bagel, based in the US and responsible for US projects and many casino projects

Paratha, based in both the UK and the US, and responsible for global projects

Crouton, based in west coast US, and responsible for PG&E and CLI

Deployment Process

This section will run through how to run a project, step by step, using a master diagram. You can view the large diagram by clicking below, or follow the process through the next few pages.

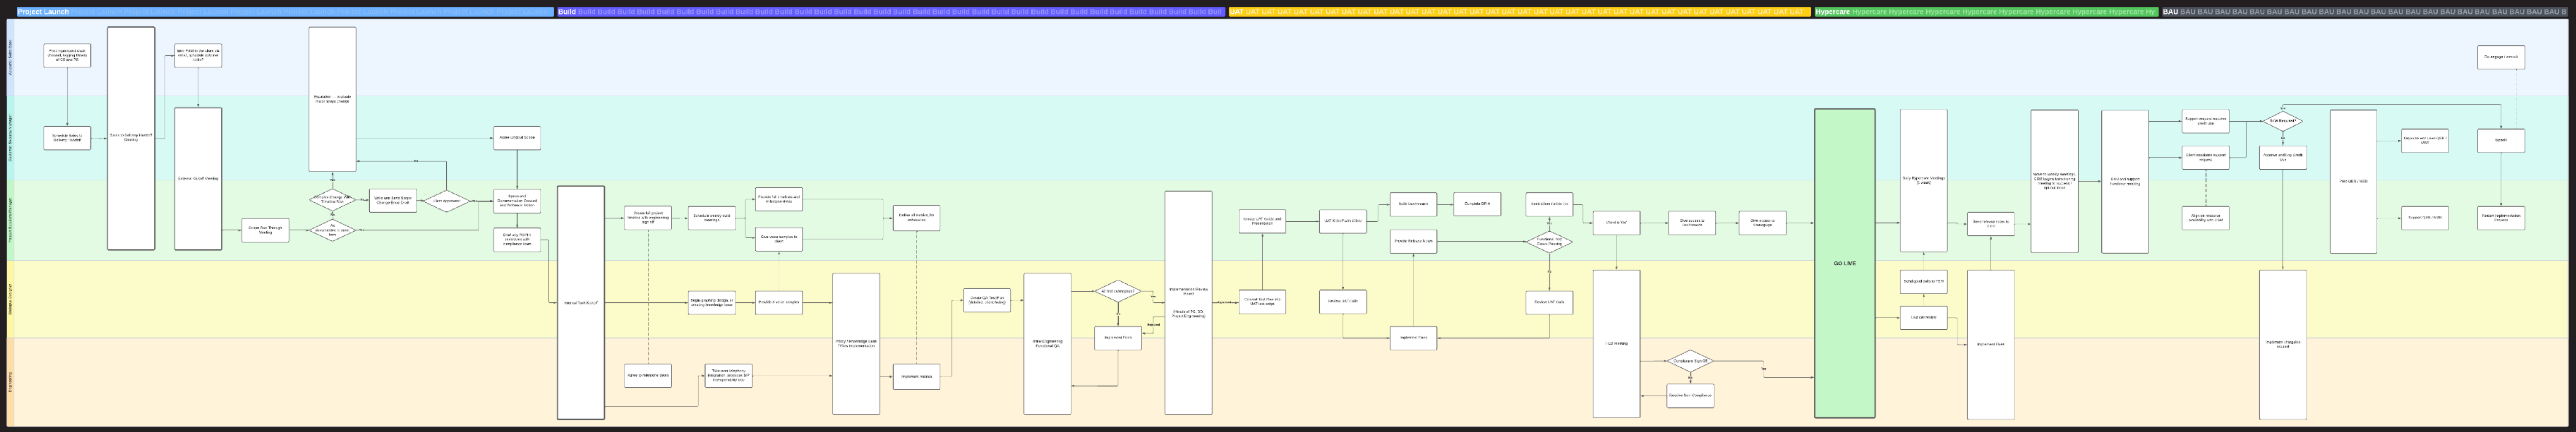

Deployment Process: Handovers

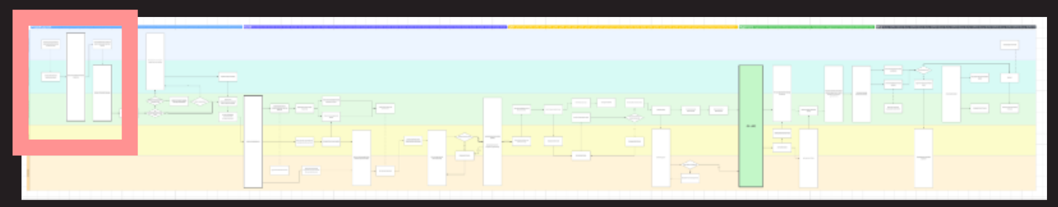

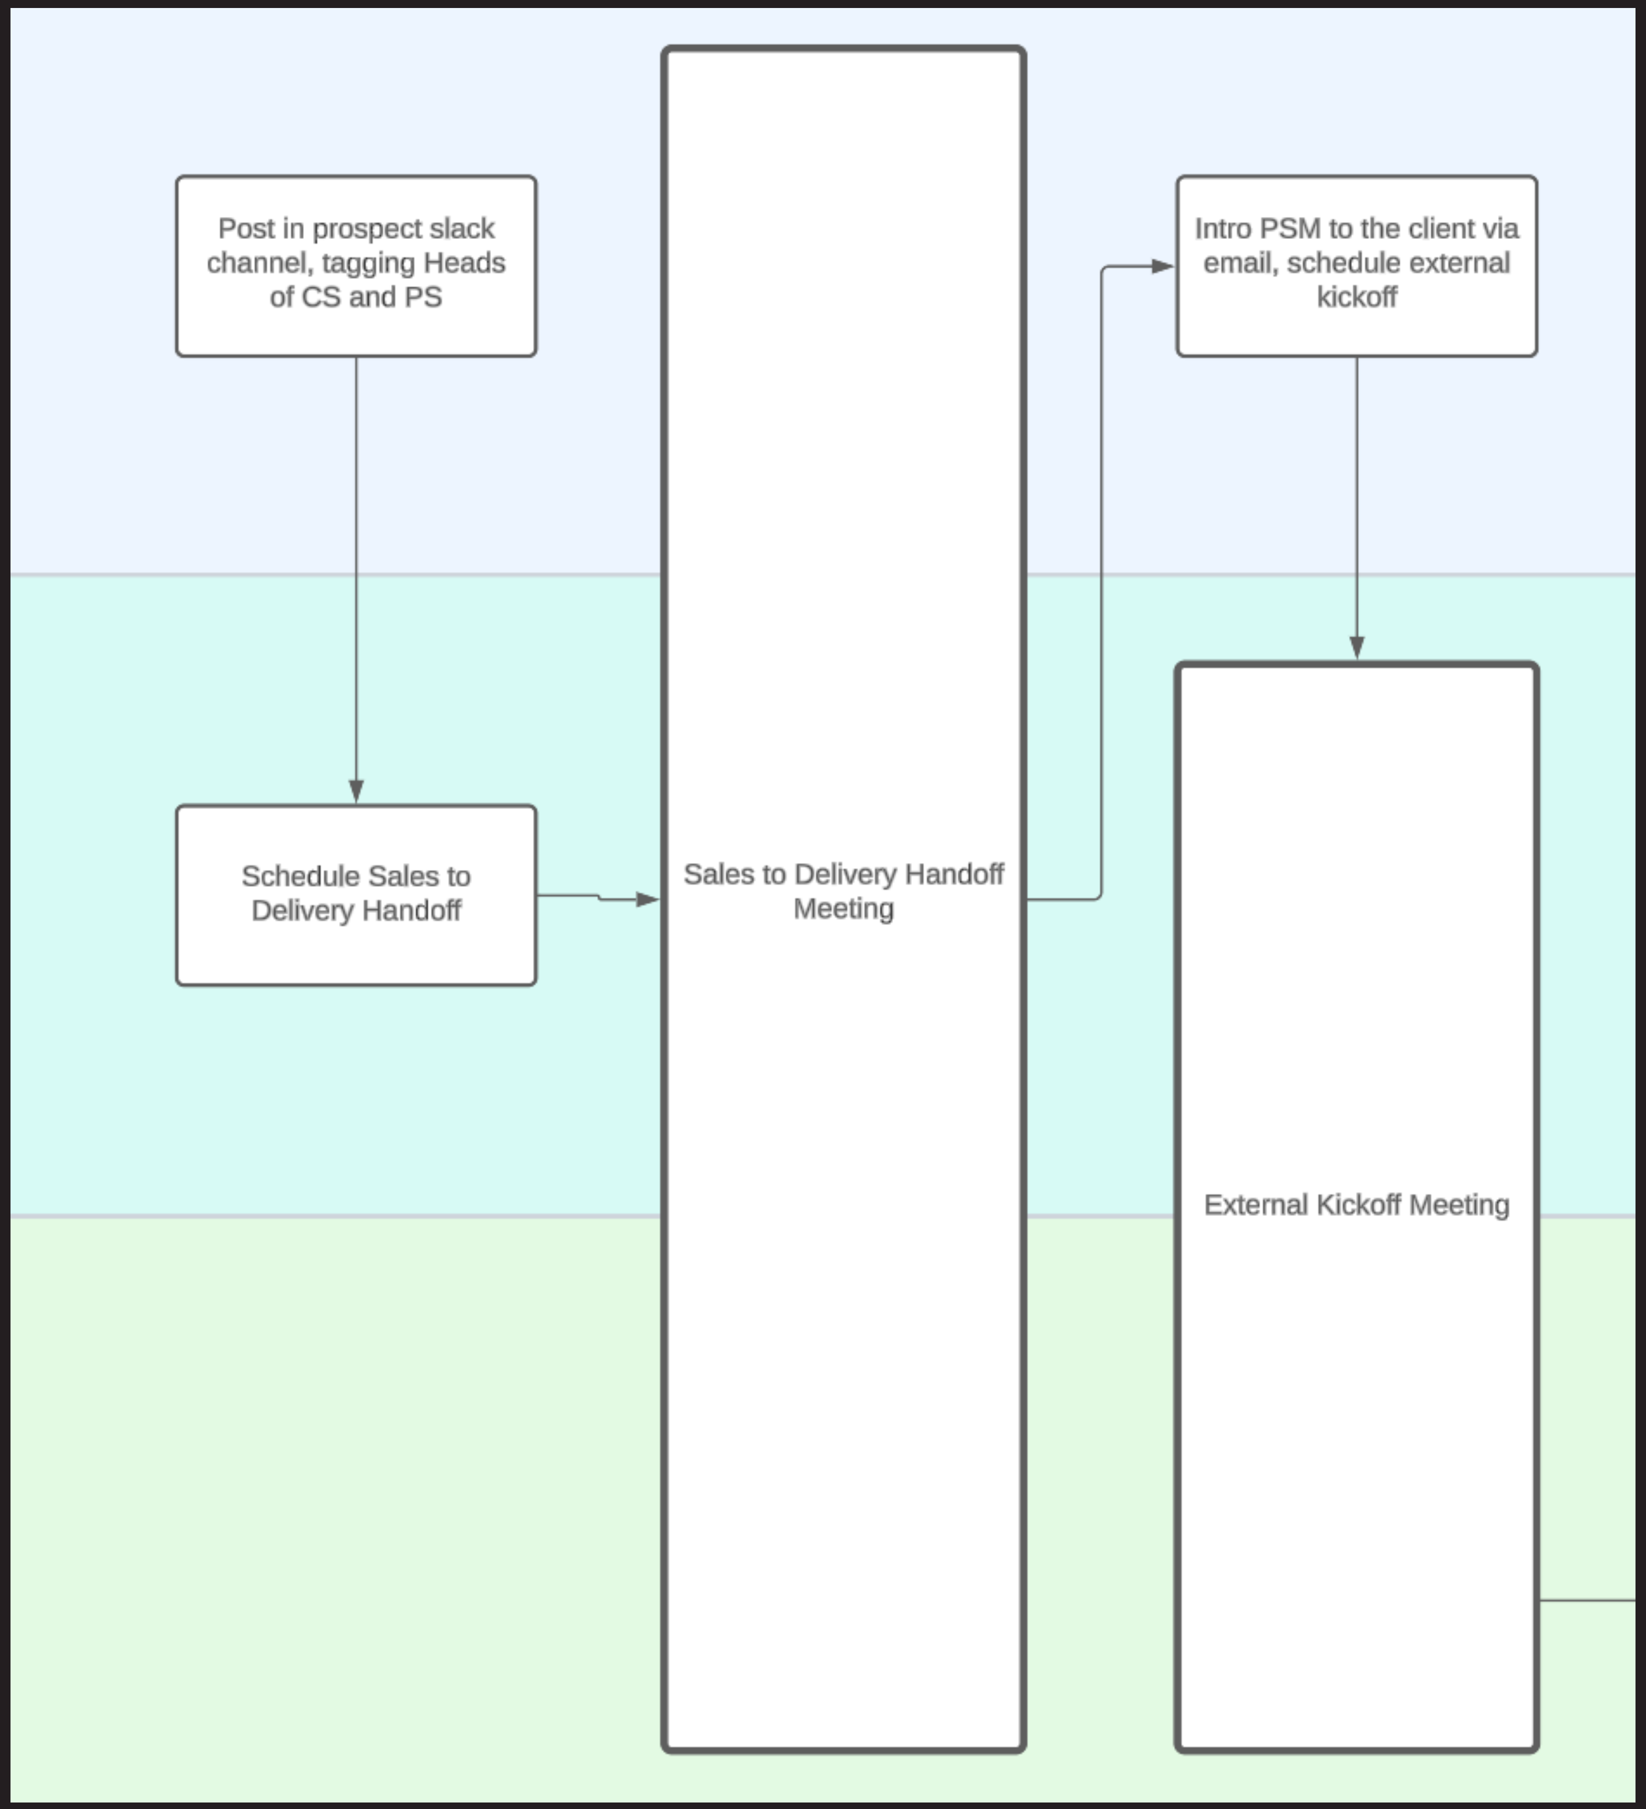

Sales to Delivery Handoff

This meeting is booked by the AE within 5-10 days of signature. At this point, a PSM is assigned to attend the meeting. Department heads will also join. This meeting has a structured slide deck. PSM's should ensure they know:

- Who the main points of contact are
- The scope, in detail, including integrations, parts that are missing, etc
- Expected timelines
- The commercial structure, ramp period, opt outs, etc

The External Kickoff Meeting

This meeting is booked by the PSM after the AE sends an introduction email. Usually, a client will invite all stakeholders to this meeting, as they expect to know who is responsible for each step of the process.

Details on how to run this meeting are found on page 28.

Deployment Process: Scope ①

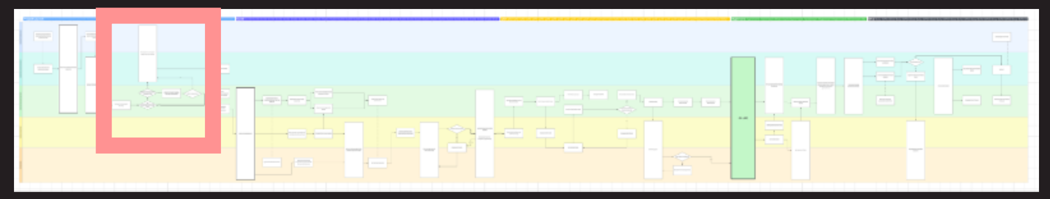

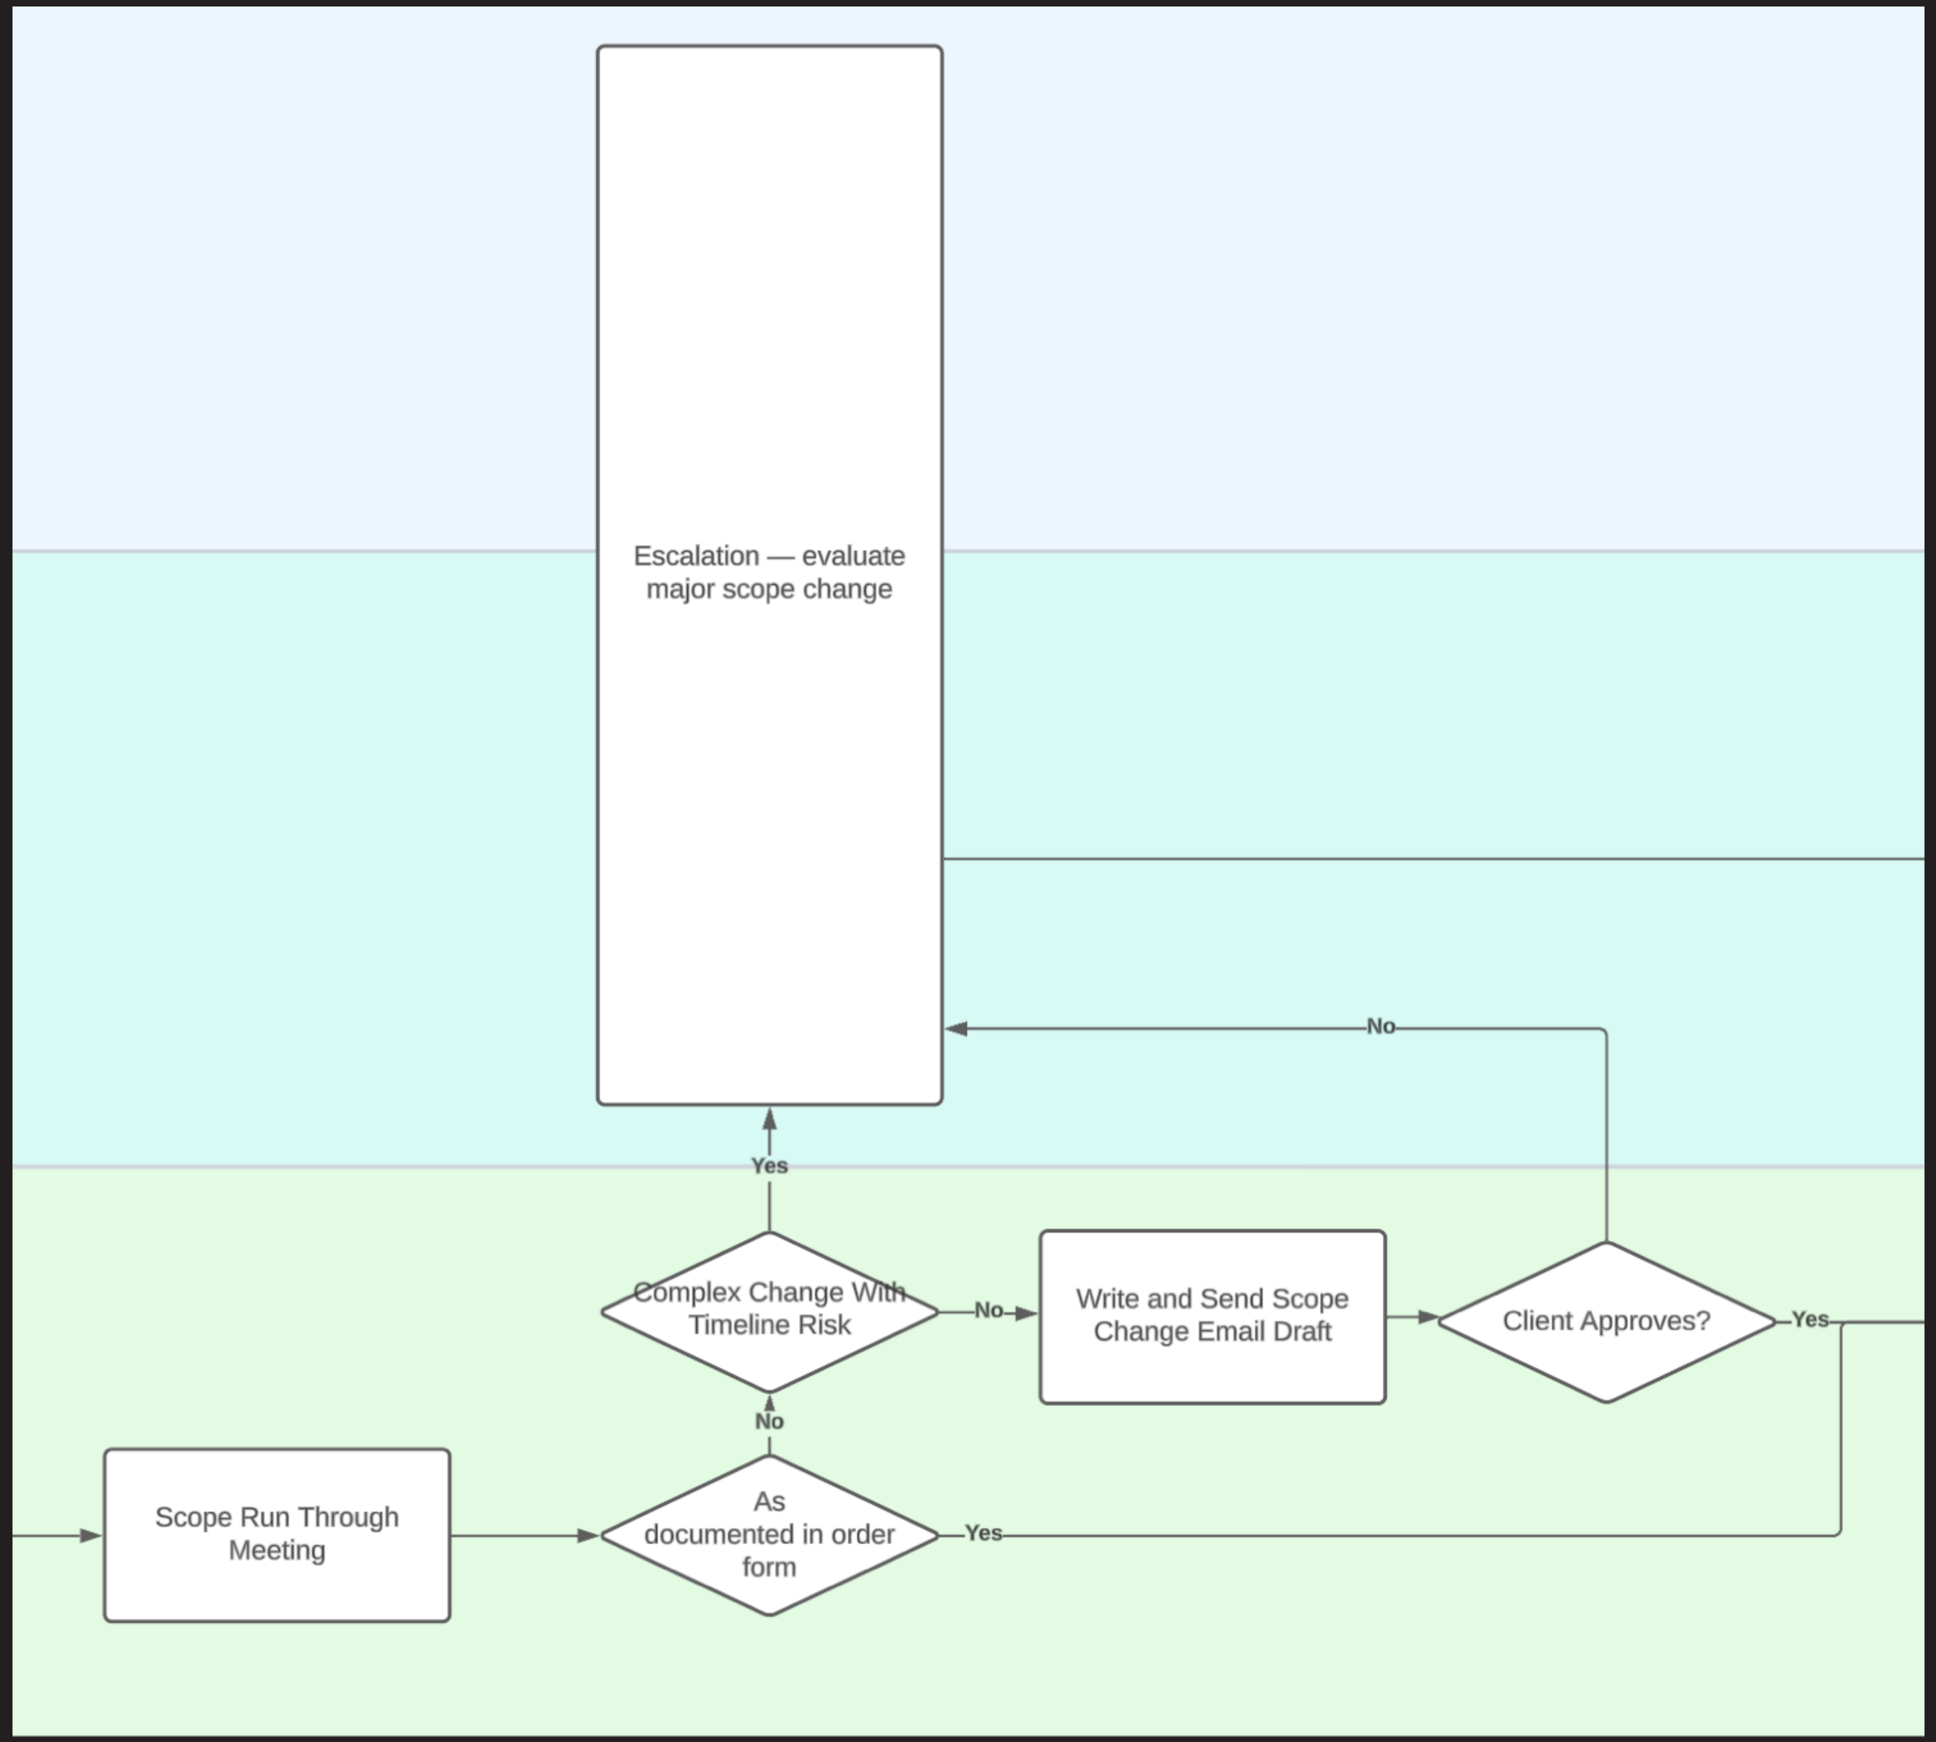

Scope Run Through

After kickoff, you should validate the scope of the virtual agent with the client. Often, the contracted scope will have nuances that will change. We want to be easy to work with, so some changes should be accepted for free.

For this meeting you should:

- Be extremely familiar with the flow in the order form. A detail oriented PSM may re-create the diagram in a more client friendly way, slightly simplified.
- Not allow the client to leave too much in the open, you want answers, and if the client cannot answer them, this may delay build
- Be able to add sample API requests and responses, and after the meeting validate all of them in the diagram
- Get final written sign off over email that the scope is final.

This is a technical call and you should feel confident that what you create on paper can go to our engineering team, and be understood.

Deployment Process: Scope ②

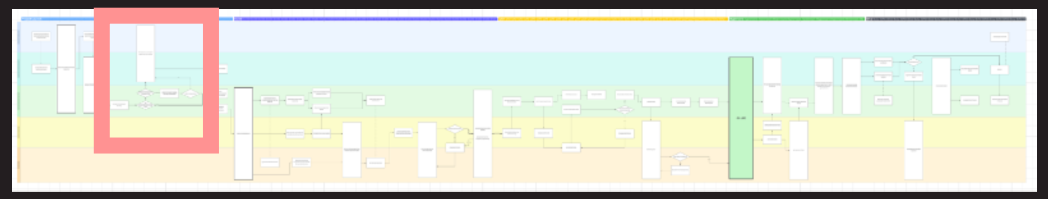

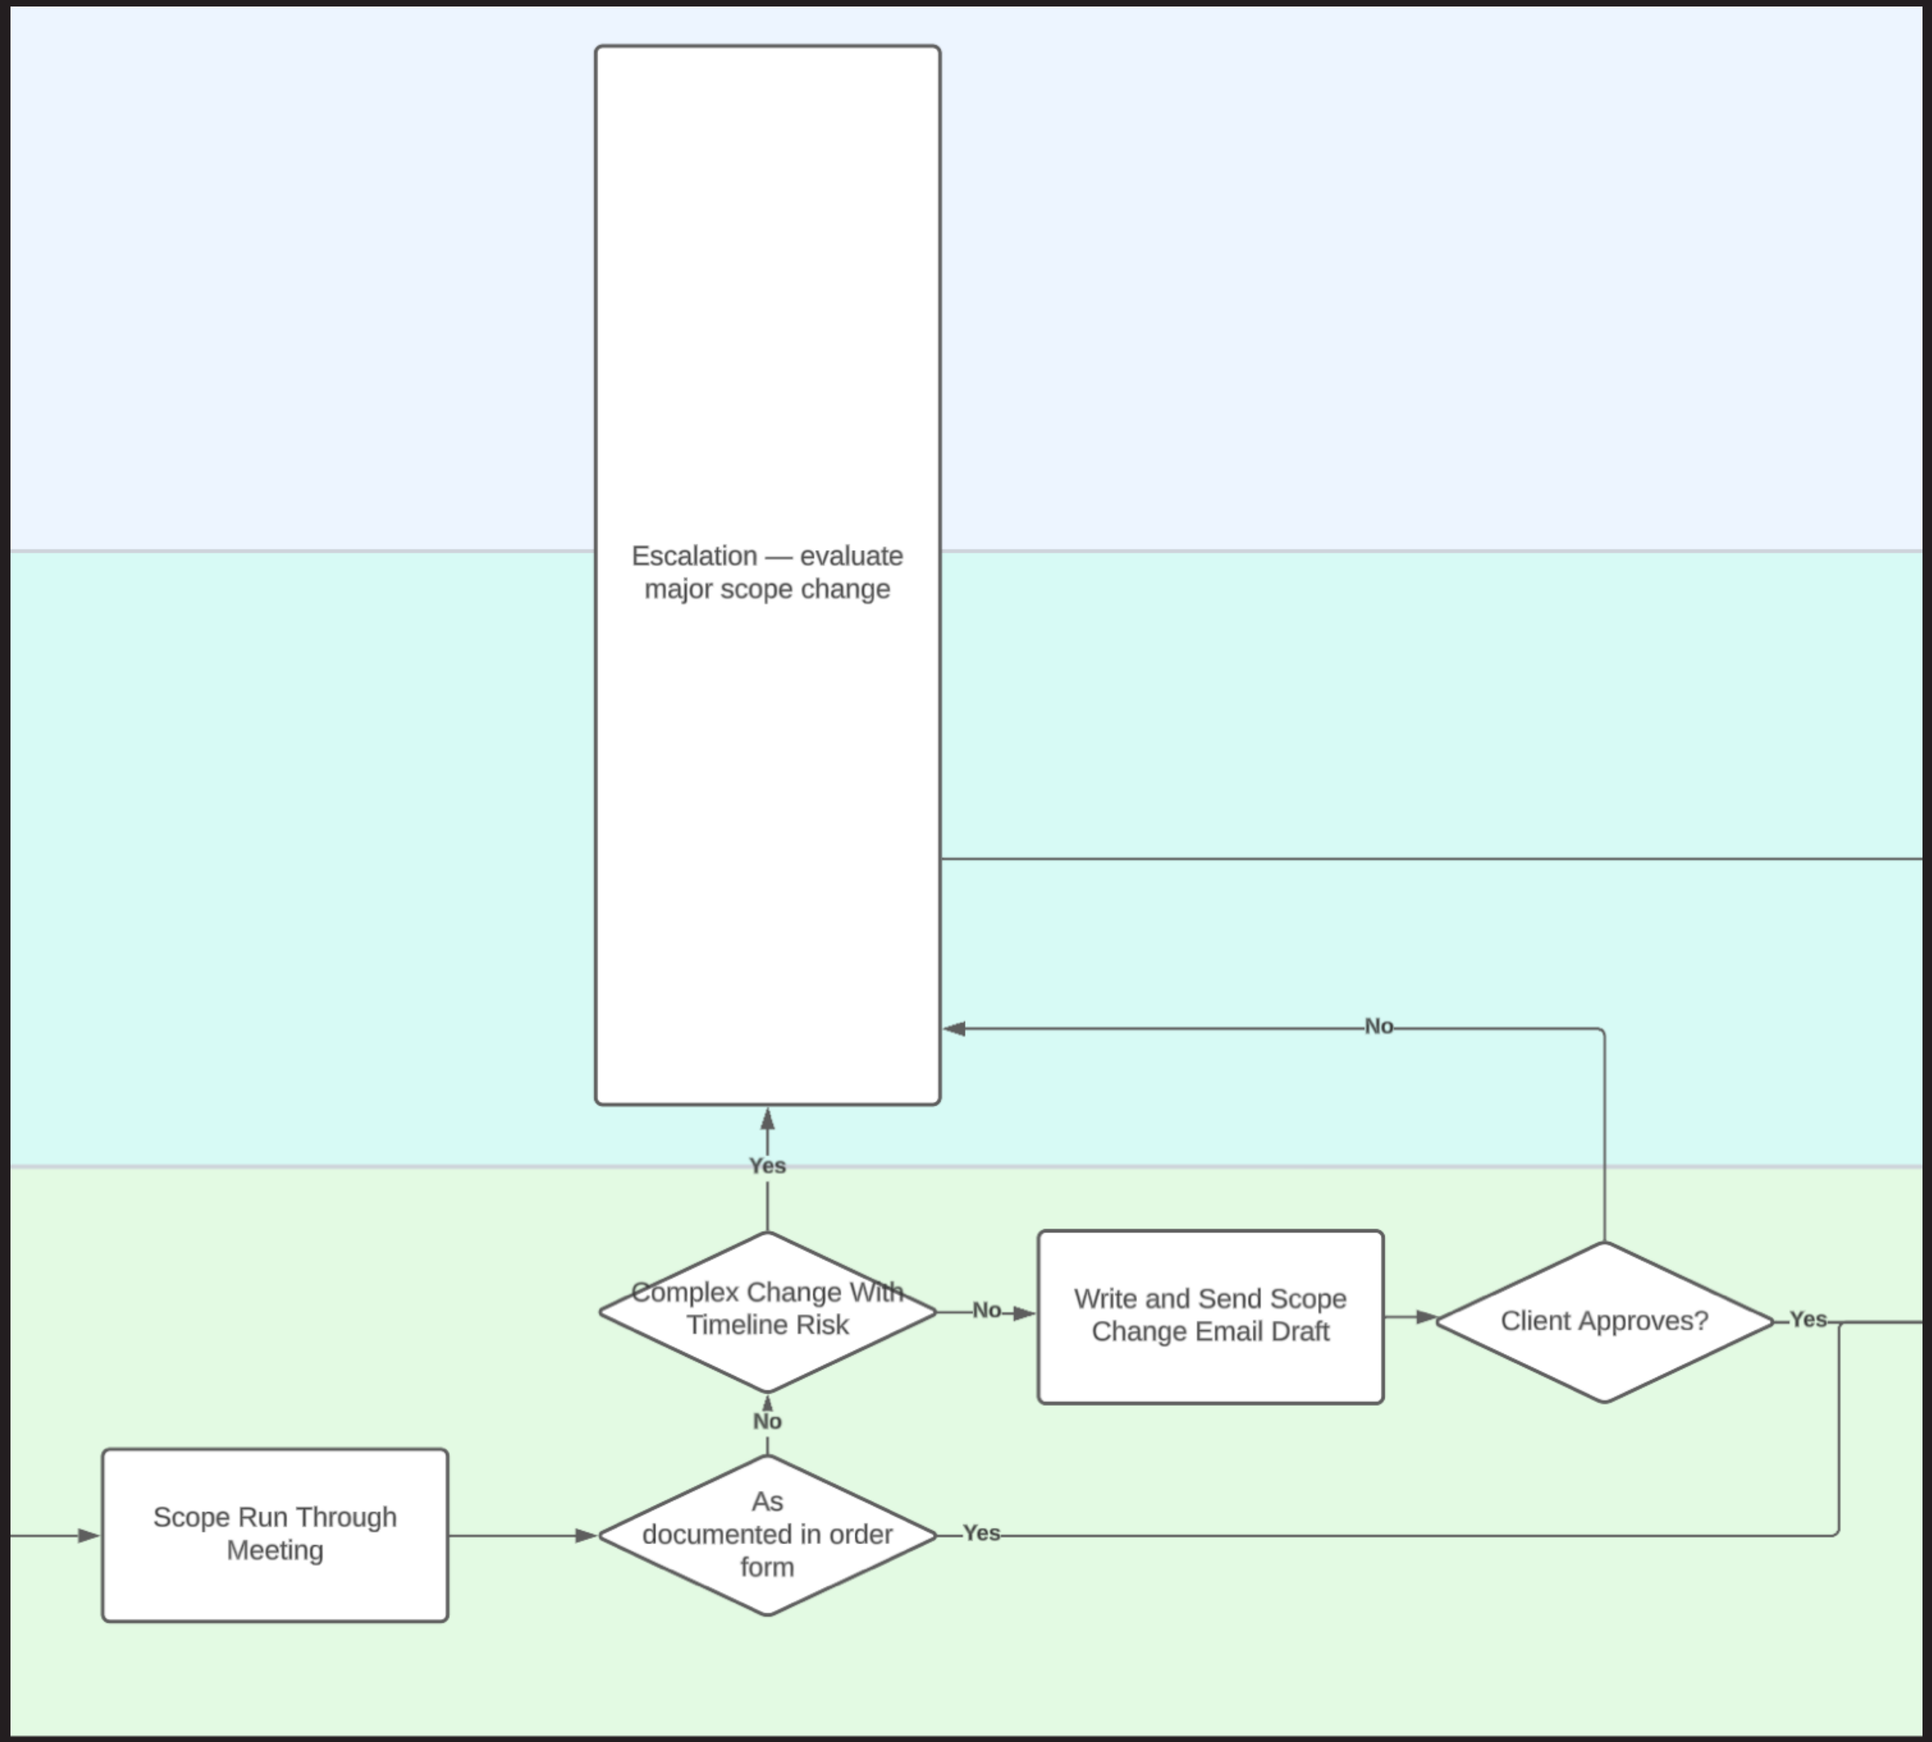

Major Scope Changes

There may be instances where clients want to change the scope substantially. If the change would require any timeline change, or significantly more effort, there are two options.

1. With the CSM, draft a Scope Change Email, outlining the change. This is a formal scope change that the client must confirm. It may come at extra cost, but this is optional. It may also invalidate opt-out, or cause a timeline extension. Because of this, you should also check with legal before sending.
2. Escalate through the AE. If the scope of the project is entirely different to what was contracted, the project needs to be transferred back to the solutions consultant team, through the AE. Send a message on the commercial slack channel, and organise a meeting with relevant team members and your line manager.

Deployment Process: Tech Kickoff

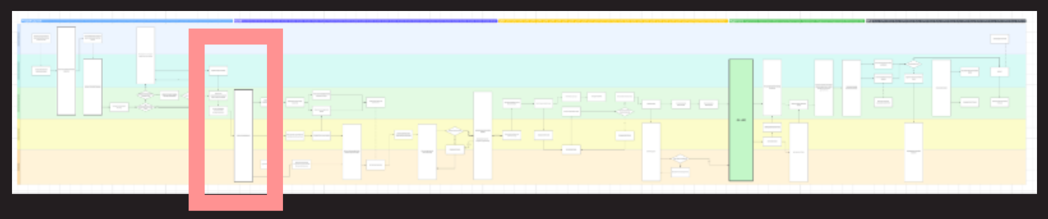

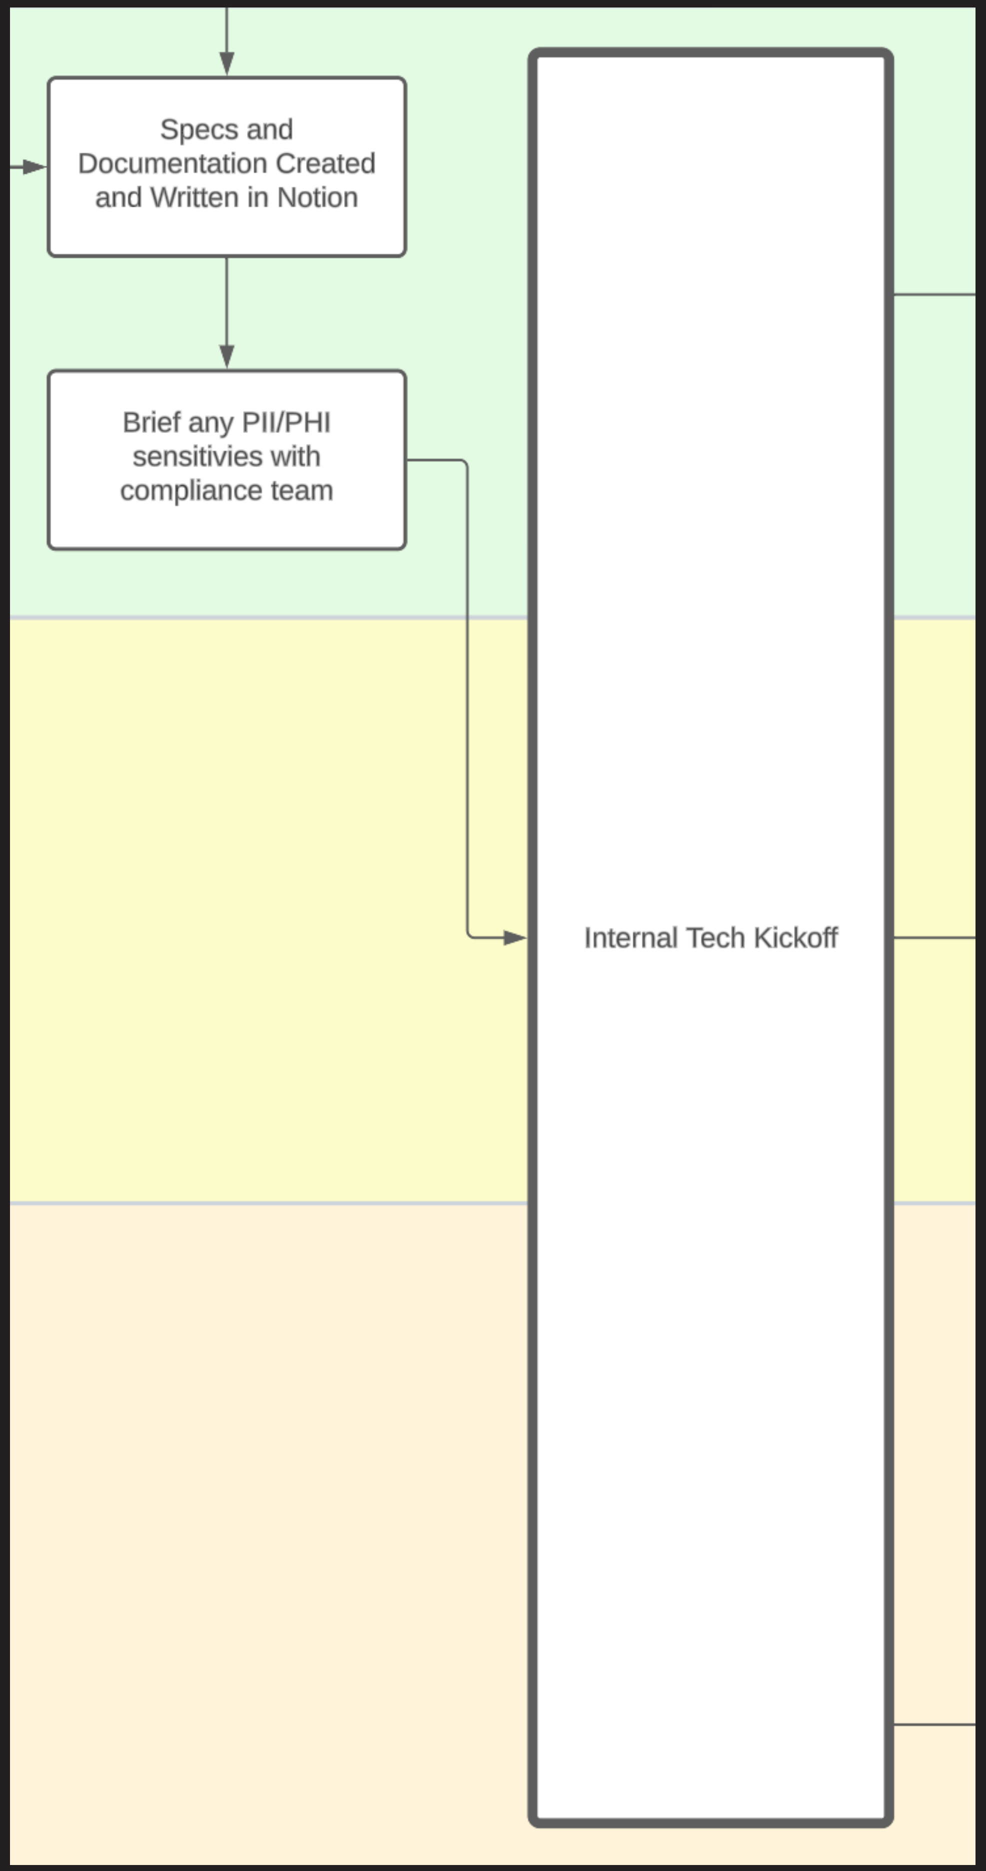

Project Setup Checklist, For the PSM

- Create a Notion project page in the projects database. Use the template
- Input contracted milestone dates into the Notion
- Place links to all discovery and contractual information into the Notion
- Create a folder for Smartsheet items in the Project Timelines workspace
- Work with the scrum master to set up a Jira board and workflow
- Create _tech and _telephony slack channels

Compliance

It is possible there will be surprises for our compliance team, especially if the project collects PII (Personal Identifiable Information) or PHI (Personal Healthcare Information). It is often good to check with our compliance team, especially if it's a new industry or use case

Internal Tech Kickoff

By this point, a pod should be fully assigned, and the PSM organises an internal tech kickoff where they run through the agreed scope, and agreeing milestone dates. At this point, the engineering lead should begin planning how long each action might take, and working backwards from the contracted go live date.

Deployment Process: Timeline

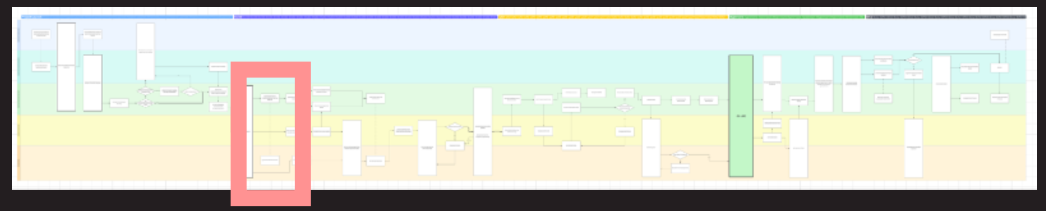

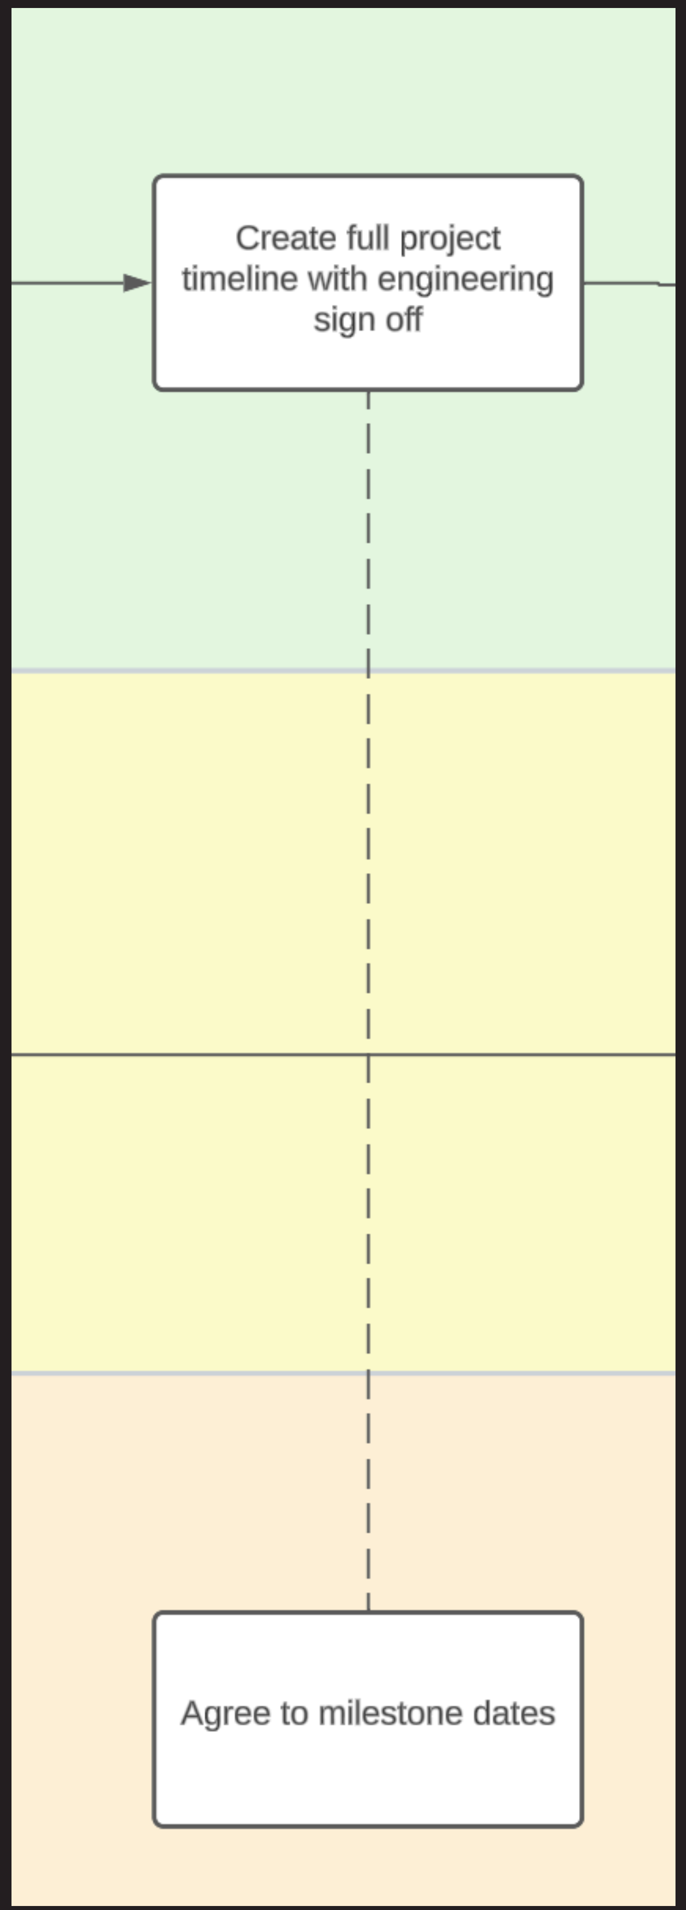

Create the Client Facing Timeline

You will create a client facing timeline, and get sign off from engineering. There is a template available in Smartsheet to use as a basis. It is important that this is up to date, day by day, and used to show the client when and how we are on track. The most effective way to hold a client to account is to assign then actions on this timeline.

You should add line items for every detailed action. You can use the lucid chart for this section of the playbook to help map out each line item. Be aware you should be able to explain them.

Work with the lead engineer in the pod on the dates.

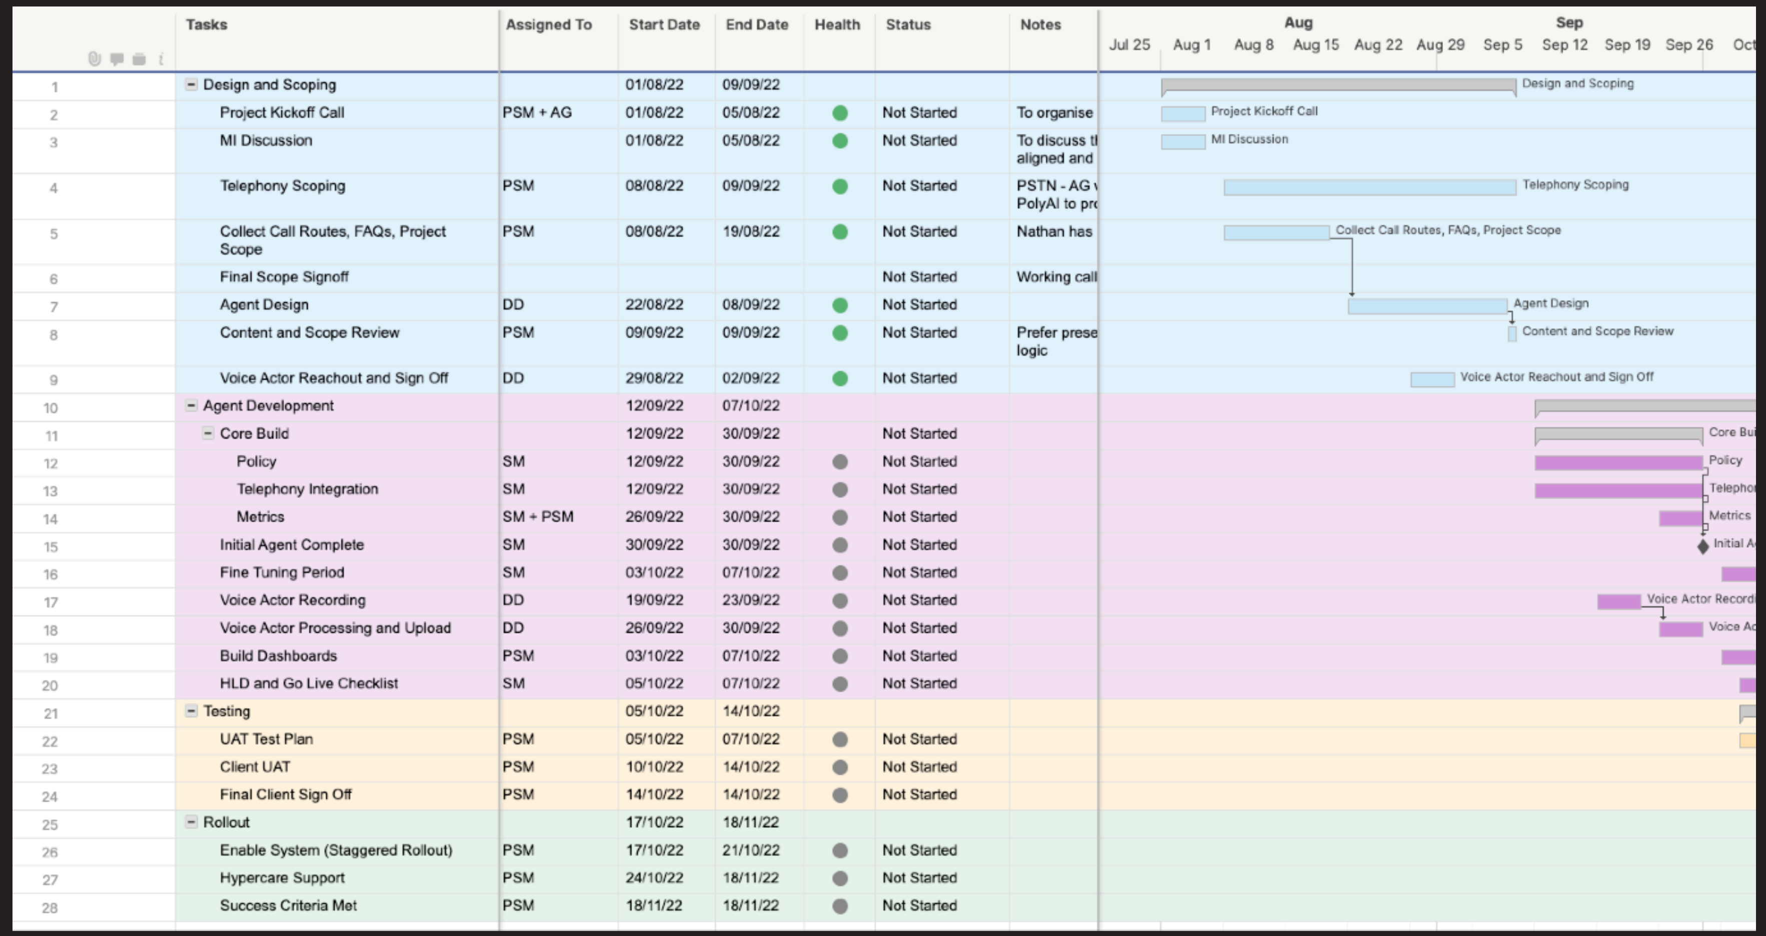

	Tasks	Assigned To	Start Date	End Date	Health	Status	Notes
1	Design and Scoping		01/08/22	09/09/22			
2	Project Kickoff Call	PSM + AG	01/08/22	05/08/22		Not Started	To organise
3	MI Discussion		01/08/22	05/08/22		Not Started	To discuss t aligned and
4	Telephony Scoping	PSM	08/08/22	09/09/22		Not Started	PSTN - AG PolyAI to pr
5	Collect Call Routes, FAQs, Project Scope	PSM	08/08/22	19/08/22		Not Started	Nathan has
6	Final Scope Signoff					Not Started	Working call
7	Agent Design	DD	22/08/22	08/09/22		Not Started	
8	Content and Scope Review	PSM	09/09/22	09/09/22		Not Started	Prefer prese logic
9	Voice Actor Reachout and Sign Off	DD	29/08/22	02/09/22		Not Started	
10	Agent Development		12/09/22	07/10/22			
11	Core Build		12/09/22	30/09/22		Not Started	
12	Policy	SM	12/09/22	30/09/22		Not Started	
13	Telephony Integration	SM	12/09/22	30/09/22		Not Started	
14	Metrics	SM + PSM	26/09/22	30/09/22		Not Started	
15	Initial Agent Complete	SM	30/09/22	30/09/22		Not Started	
16	Fine Tuning Period	SM	03/10/22	07/10/22		Not Started	
17	Voice Actor Recording	DD	19/09/22	23/09/22		Not Started	
18	Voice Actor Processing and Upload	DD	26/09/22	30/09/22		Not Started	
19	Build Dashboards	PSM	03/10/22	07/10/22		Not Started	
20	HLD and Go Live Checklist	SM	05/10/22	07/10/22		Not Started	
21	Testing		05/10/22	14/10/22			
22	UAT Test Plan	PSM	05/10/22	07/10/22		Not Started	
23	Client UAT	PSM	10/10/22	14/10/22		Not Started	
24	Final Client Sign Off	PSM	14/10/22	14/10/22		Not Started	
25	Rollout		17/10/22	18/11/22			
26	Enable System (Staggered Rollout)	PSM	17/10/22	21/10/22		Not Started	
27	Hypercare Support	PSM	24/10/22	18/11/22		Not Started	
28	Success Criteria Met	PSM	18/11/22	18/11/22		Not Started	

Using Smartsheet

Smartsheet can be tricky to use. However, as these are client facing it is imperative that they are clean, tidy, and professional. You will want to use this to hold clients to account, and if they are untidy or confusing, they will not be taken seriously.

Use the table view to enter dates, and then visualise it in the Gantt view afterwards. Do not set dependencies until you have finished setting dates, however do use the dependency feature to demonstrate to clients the cascading effect of them being late. Keep the status up to date at all times.

Use the Smartsheet help guides.

Deployment Process: Building ①

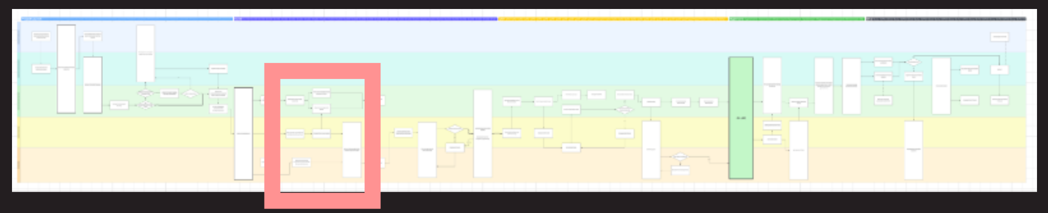

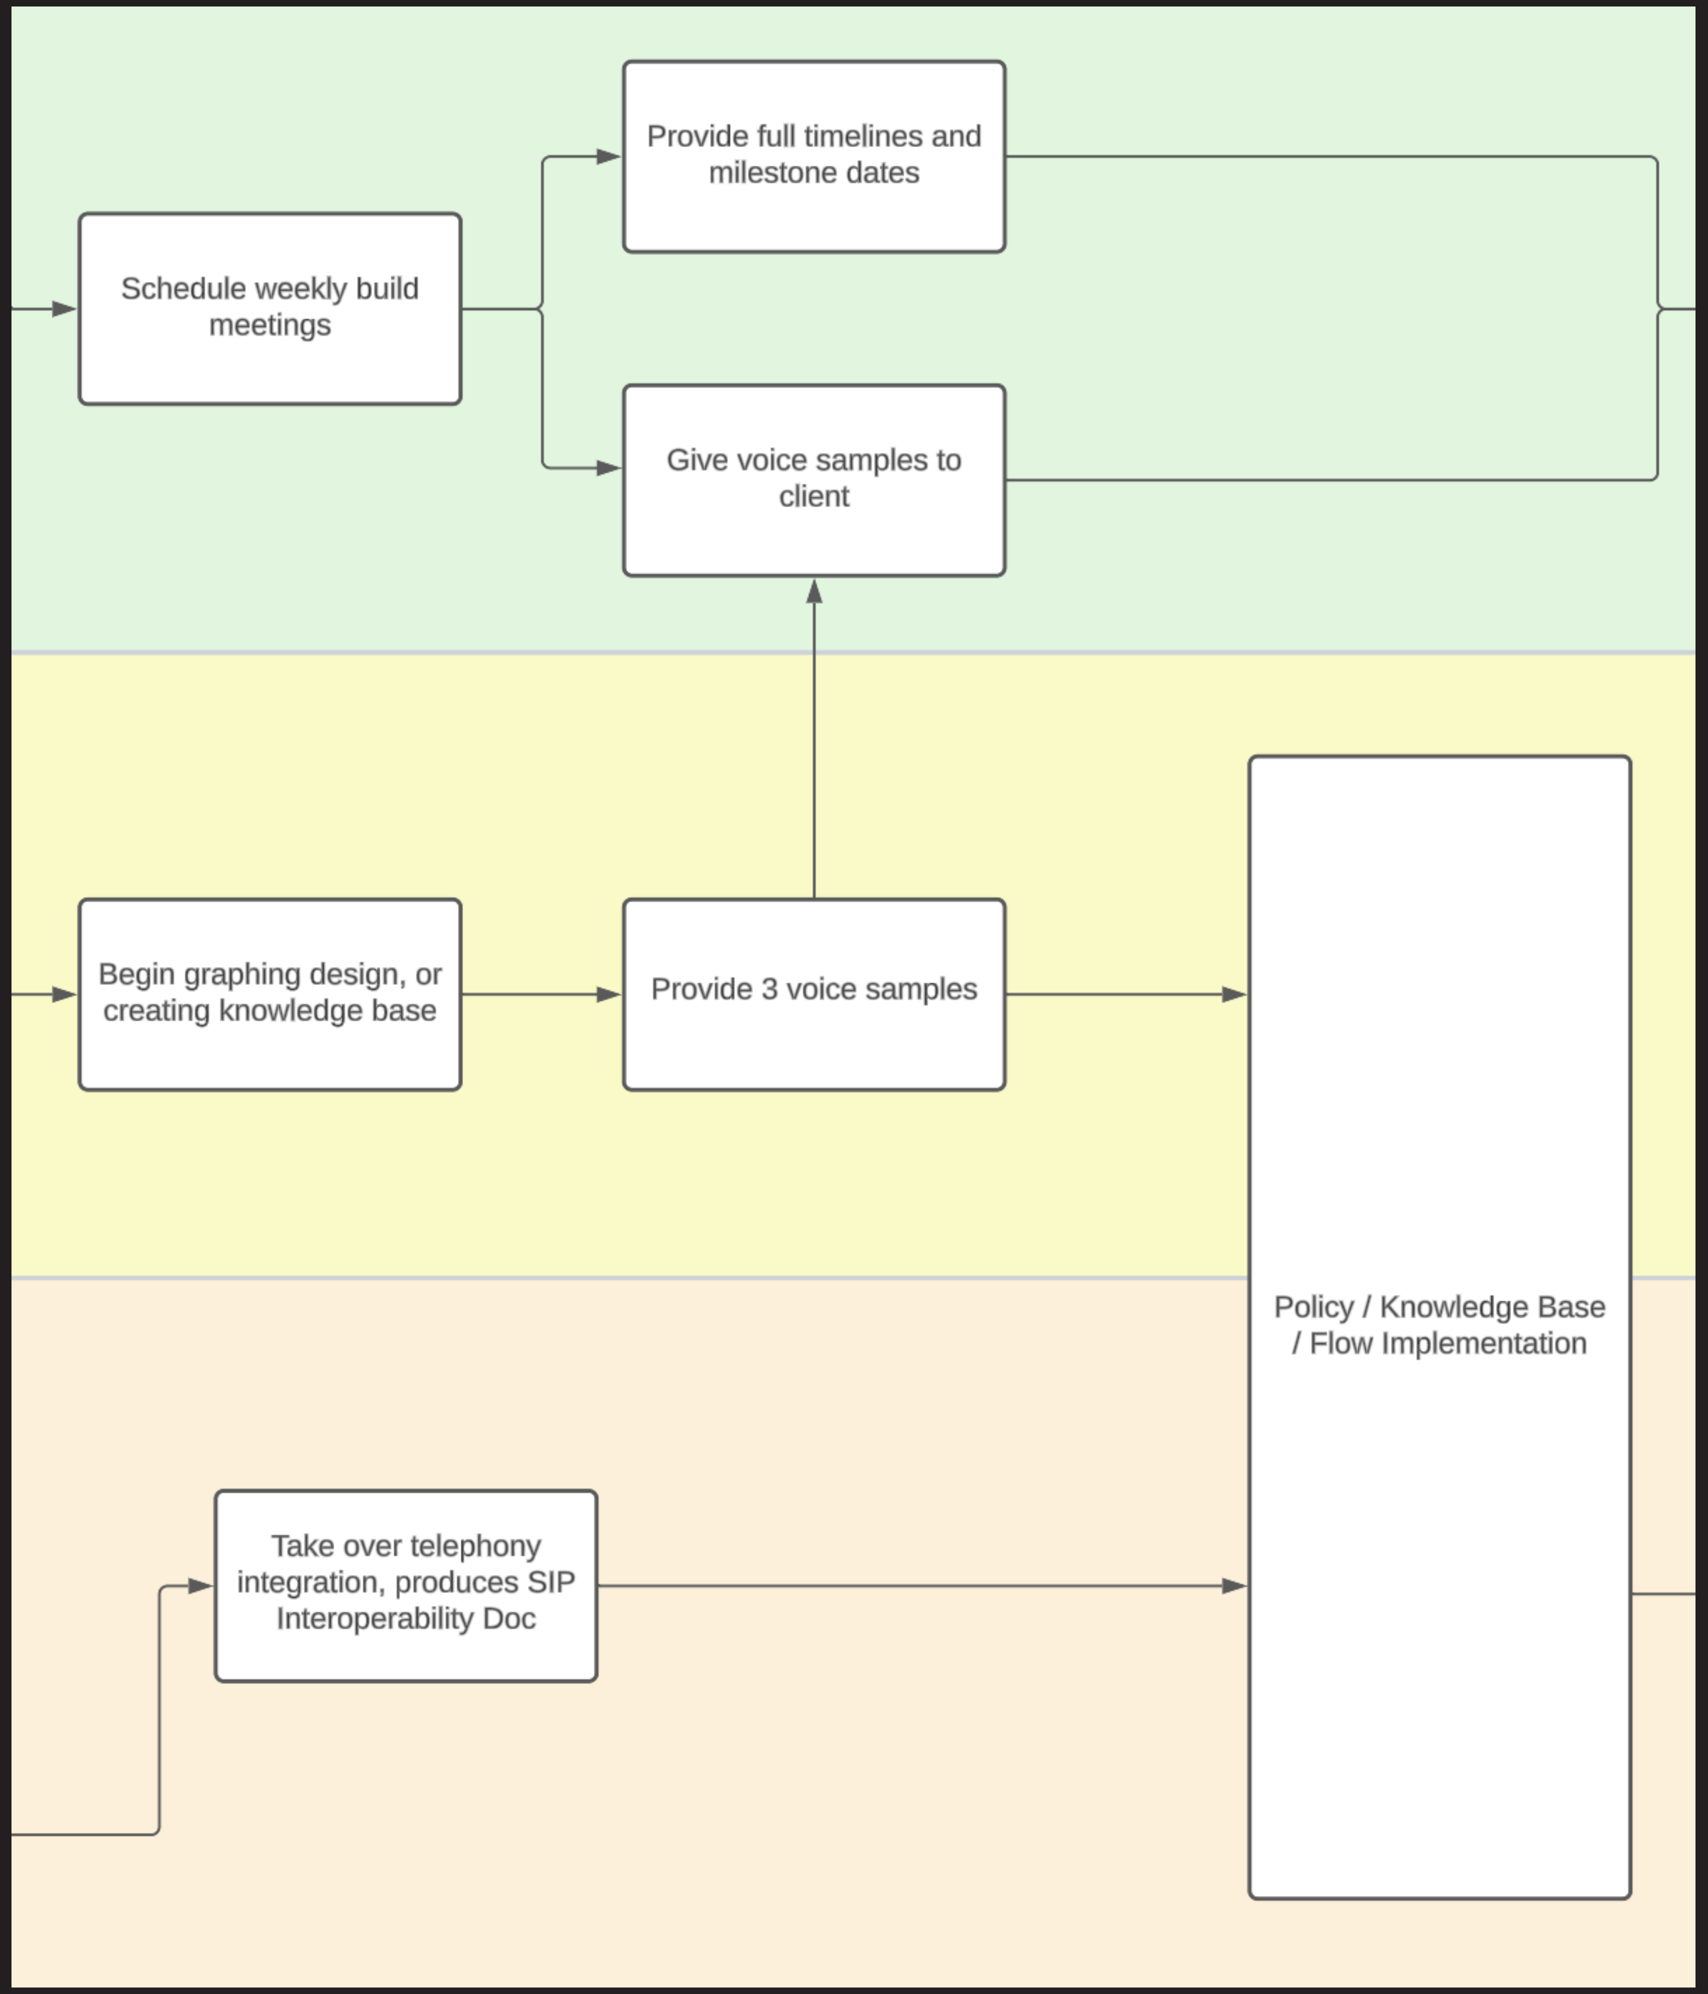

Client Meetings

We run weekly meetings with the client, often more, to keep them up to date with progress. In these meetings you will show the action tracker, the timeline and report progress. You will play voice samples to the client and coach them into picking the right one. Work closely with the Agent Designer to pick the right voice.

We don't usually ask for preferences, rather send voices that we think would work for the agent, and then take feedback.

It's also often wise to get an early recording of the real system mid-development to show progress.

More guides on how to run such daily meetings can be found on page X.

Deployment Process: Building ①

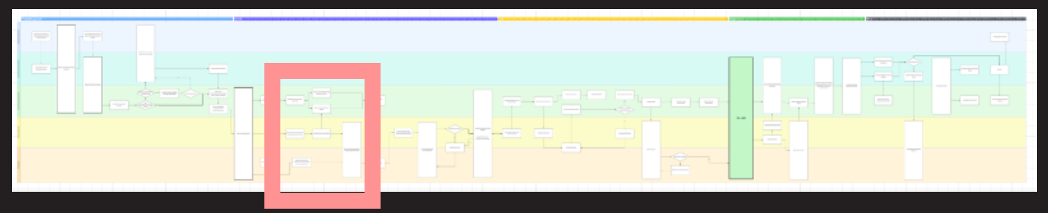

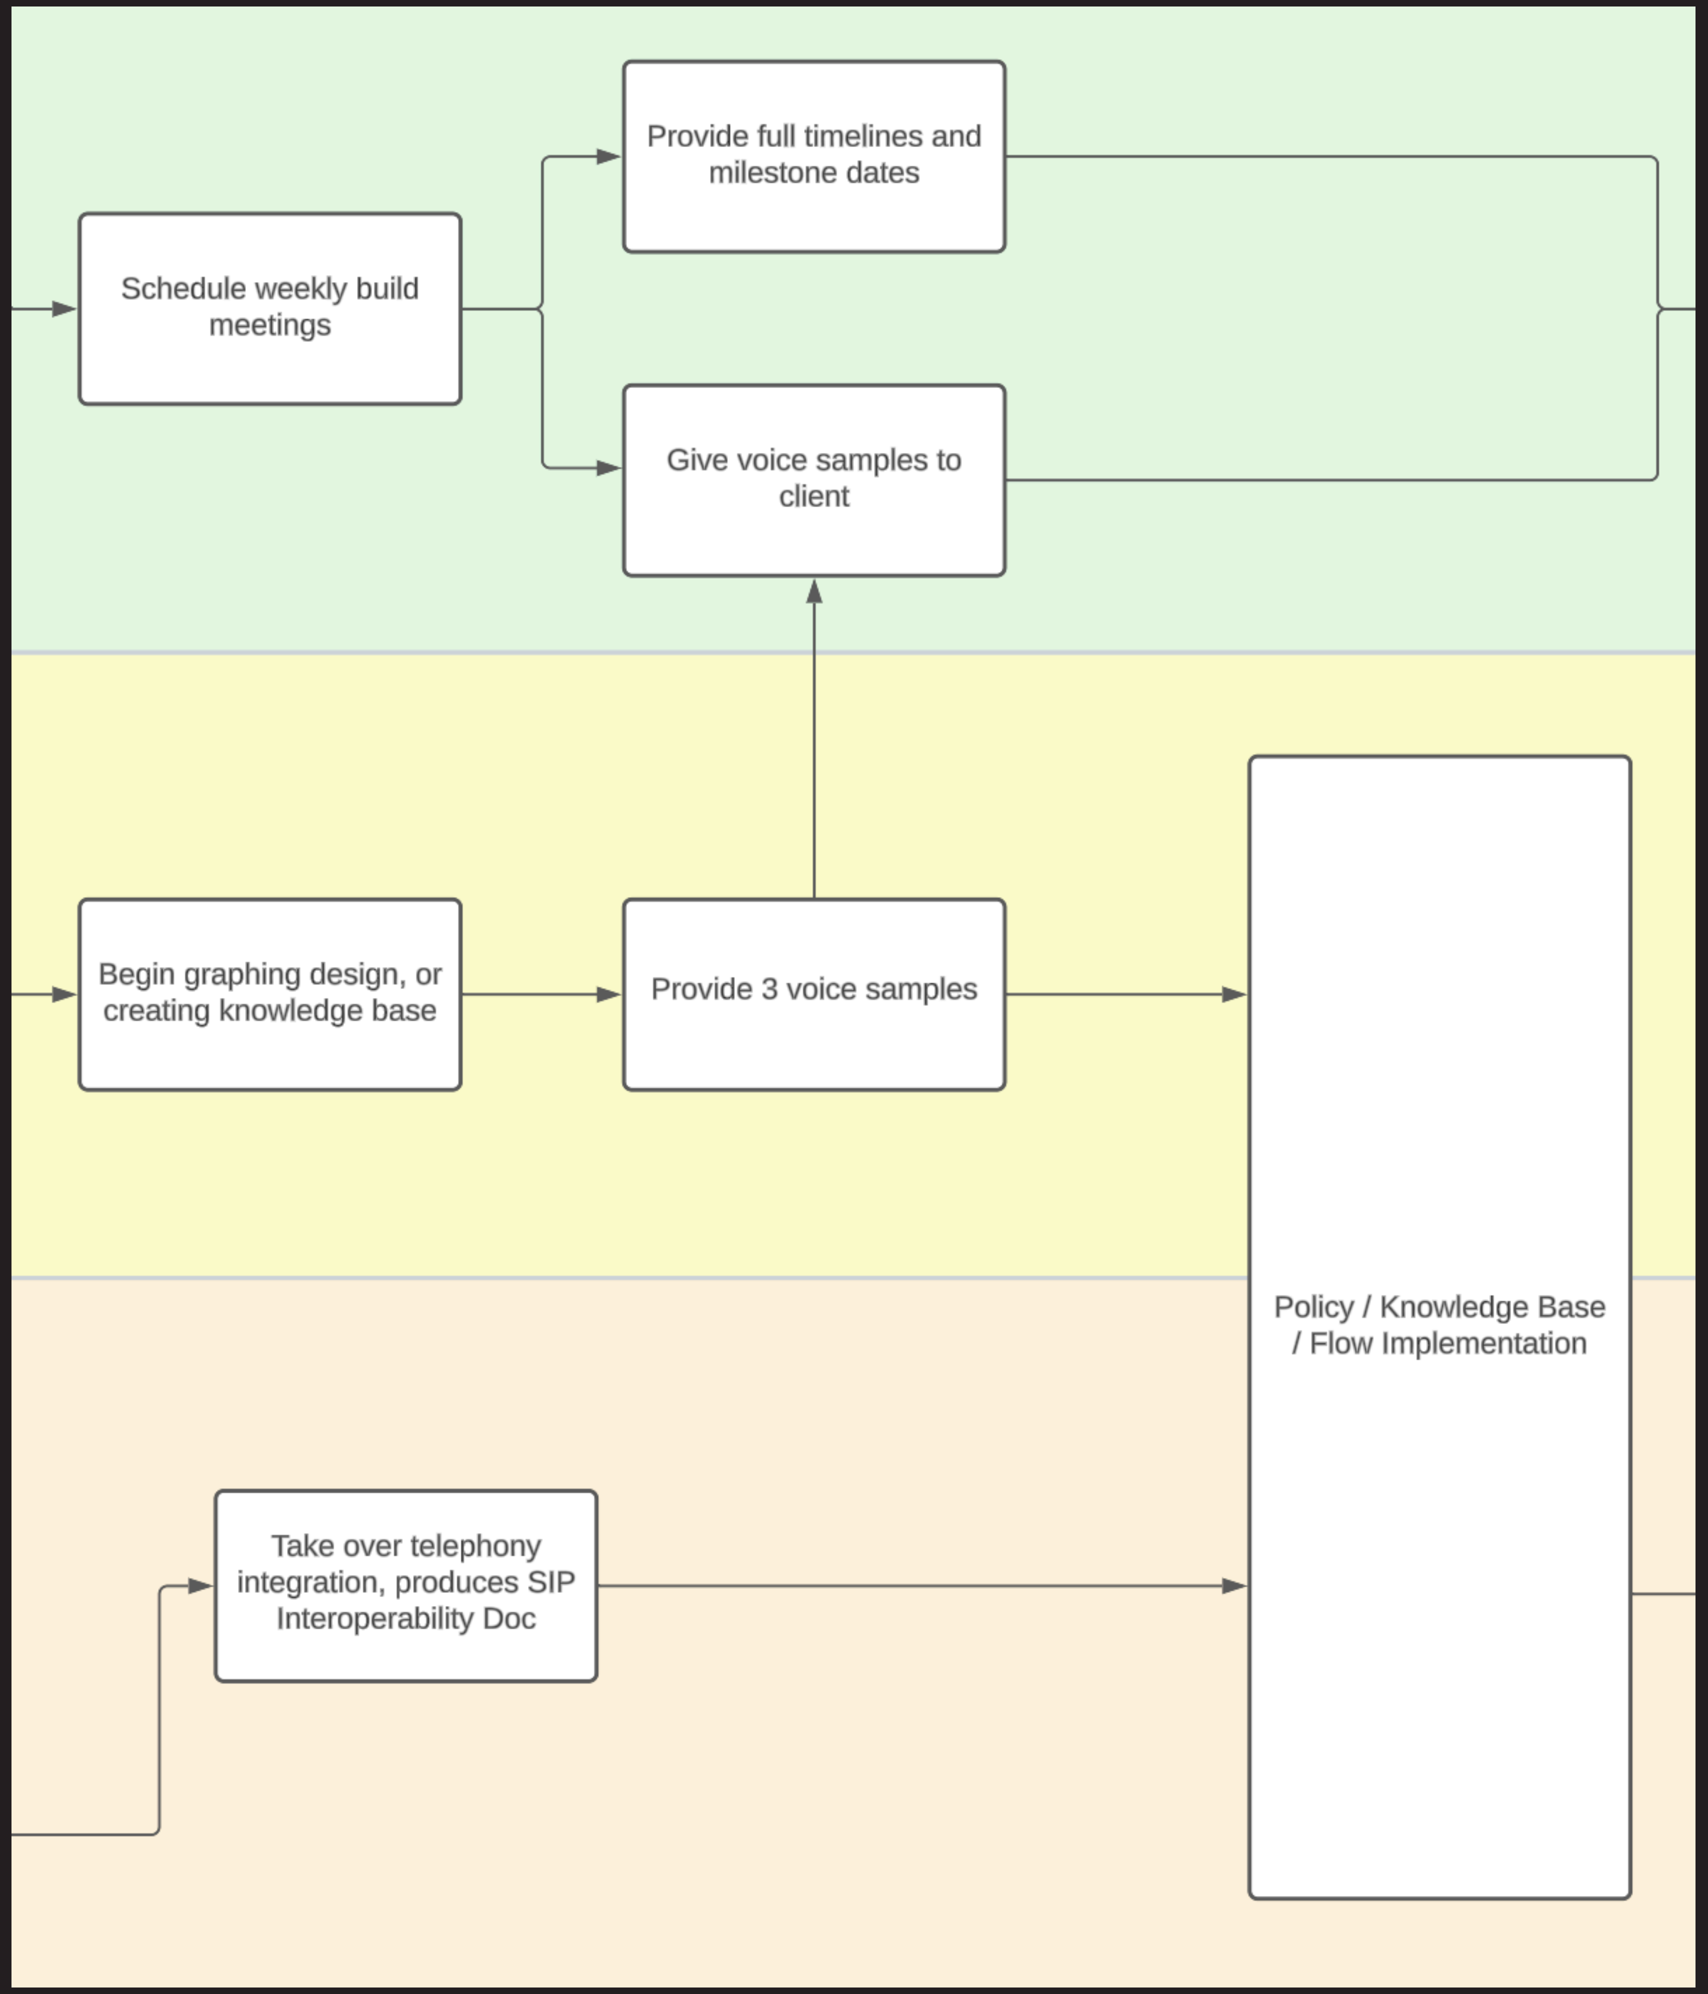

Telephony Integration and Interoperability document

Your telephony engineer will own this document, however check in frequently to ensure it is up to date, and that the client is also following their deliverables on time. It is very easy for the deliverables in this document to slip.

Providing voice samples

Your Agent Designer will generate samples. Listen to them first, and make sure they are perfect. It's a little tedious to generate samples so don't be afraid to push back if you find the voice strange.

You should provide **no more than three** samples to the client upfront, and play them on a call (there is a template here 🔗). The more you provide up front, the more limited the options if they don't like any.

You should always include 'PolyVoices' in the samples. These are voices where we own the relationship with the voice actor and so can get live recordings later for unusual pronunciations, or to maximise agent voice quality. Our clients have been sold superhuman assistants, so PolyVoices give us the most flexibility.

Deployment Process: Building ②

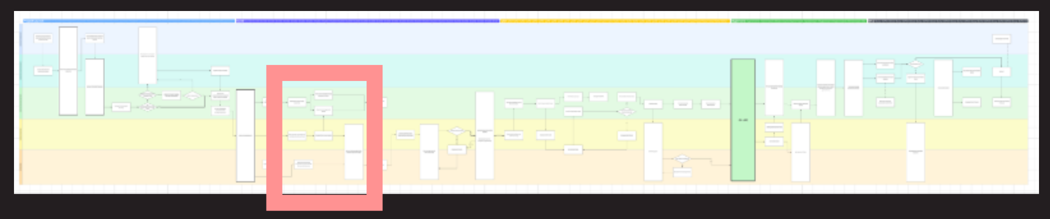

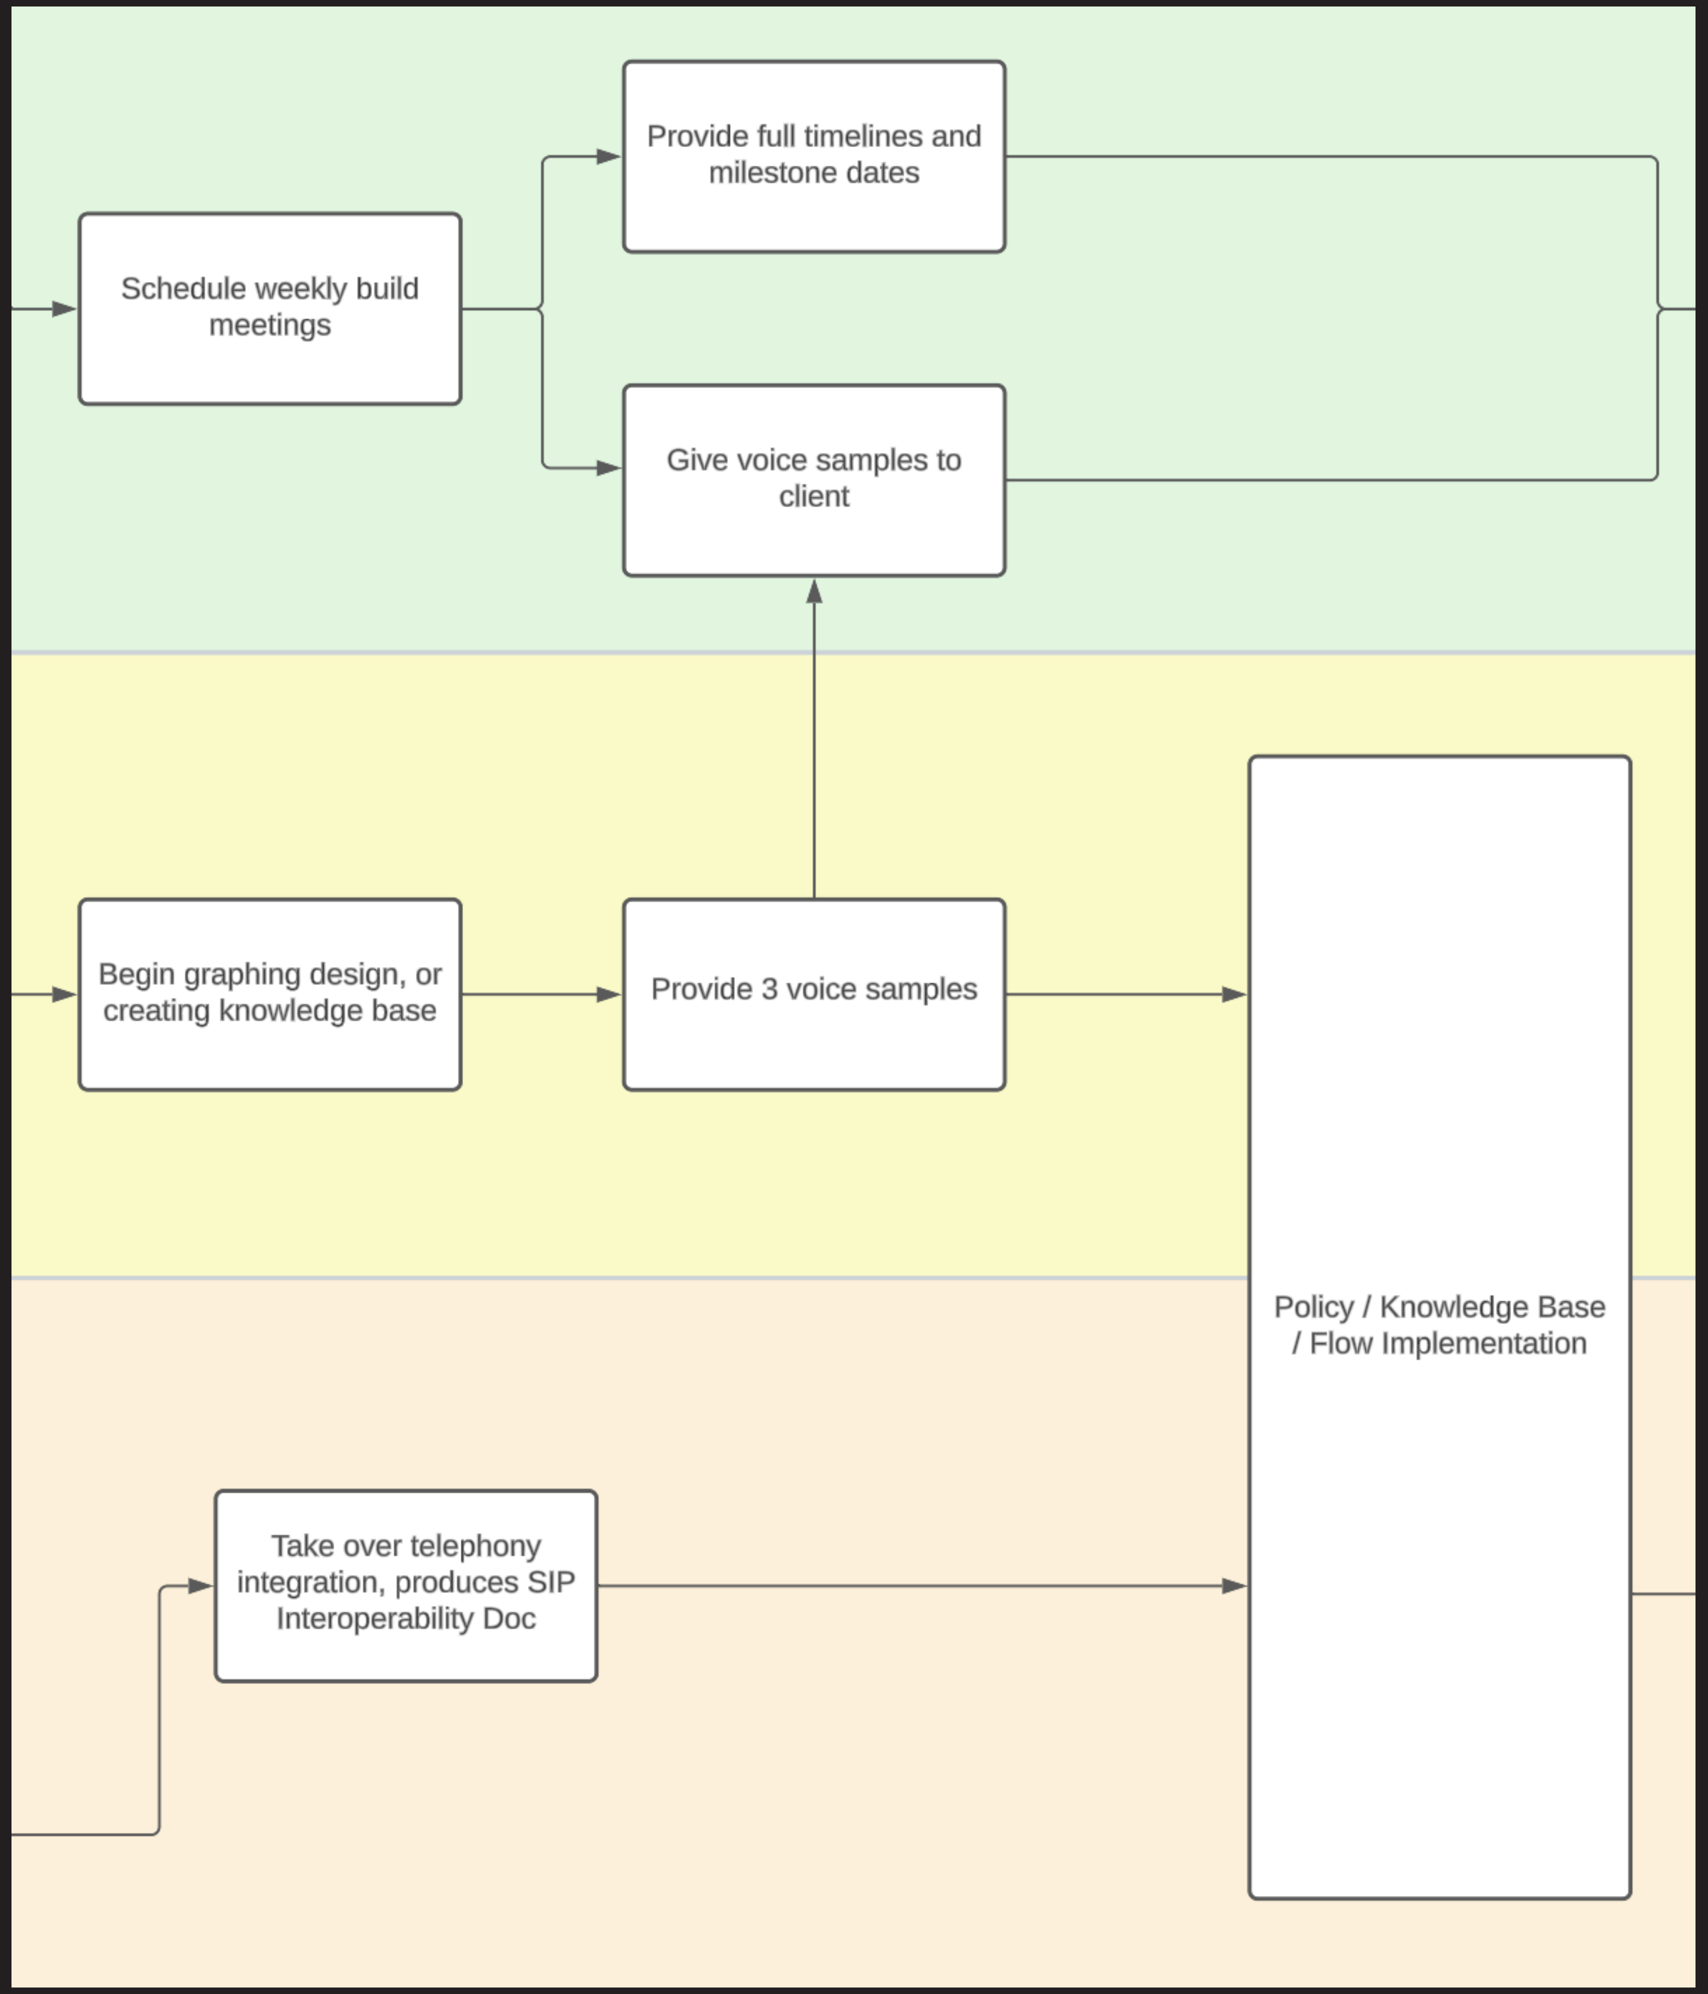

Keeping Close to the Engineering Team

WAs a PSM, the engineering team are ultimately waiting on your guidance and timelines. Projects will not move at the pace required unless you are close to the day to day. You must:

- Be deeply familiar with the bot and it's behaviour. Clients will ask you about it, and you should be able to answer forcefully and confidently. The bot's design is as much yours as it is the team's.
- Attend scrum to be aware of day to day deliverables
- Be present in the tech channel, if you are not seeing progress, there may not be progress. Nobody wants to micromanage, but you must be confident and having the team report progress. Delays, poor agent quality and client frustration will come if you cannot justify it.

Deployment Process: Metrics

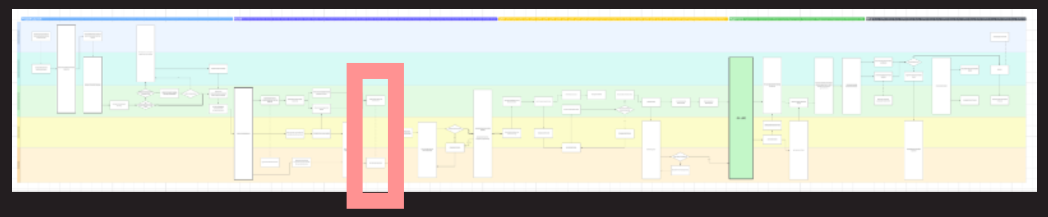

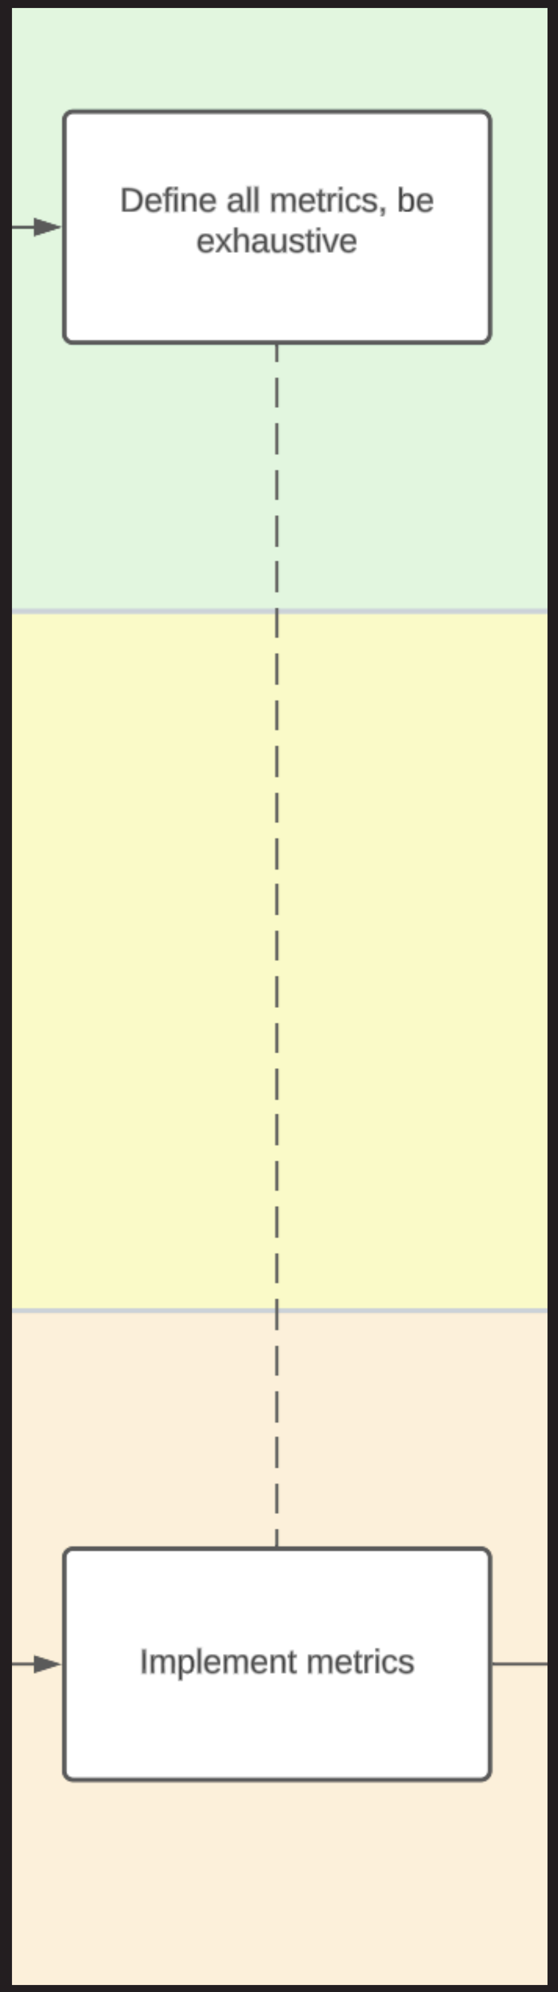

Defining Metrics

The PSM is responsible for defining the metrics required in the project. There are some standard metrics which must be followed and you should check that these are implemented correctly.

Metrics are ‘checkpoints’ in a conversation. When a caller hits a point in a conversation, a metric will be written. It could be something simple like ‘BOOKING_INITIATED’ which is a single event. Or it could contain data, for example ‘QA’ which signals a question was asked but will have a value such as ‘PARKING’, the name of the FAQ.

Detail on standardised metrics can be found on Notion. However, there is also a new effort to improve it, which will be documented soon.

The Consequences of Incorrect Metrics

Metrics are how we measure business cases, and more, with our clients. Getting it wrong can create huge problems and require time consuming backfilling and extended manual call listening.

You should extensively QA that metrics are being written. QA documents must have an ‘expected metrics’ column, and you must validate these on every test call.

Deployment Process: QA and Review

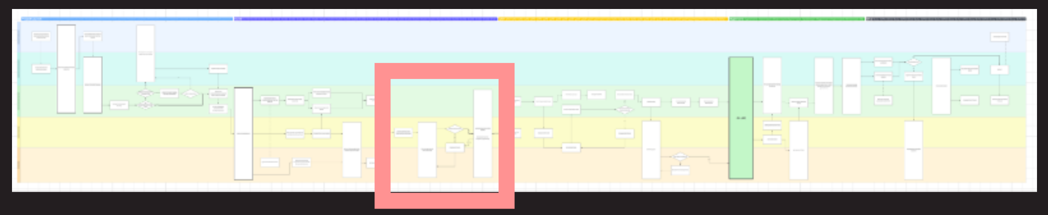

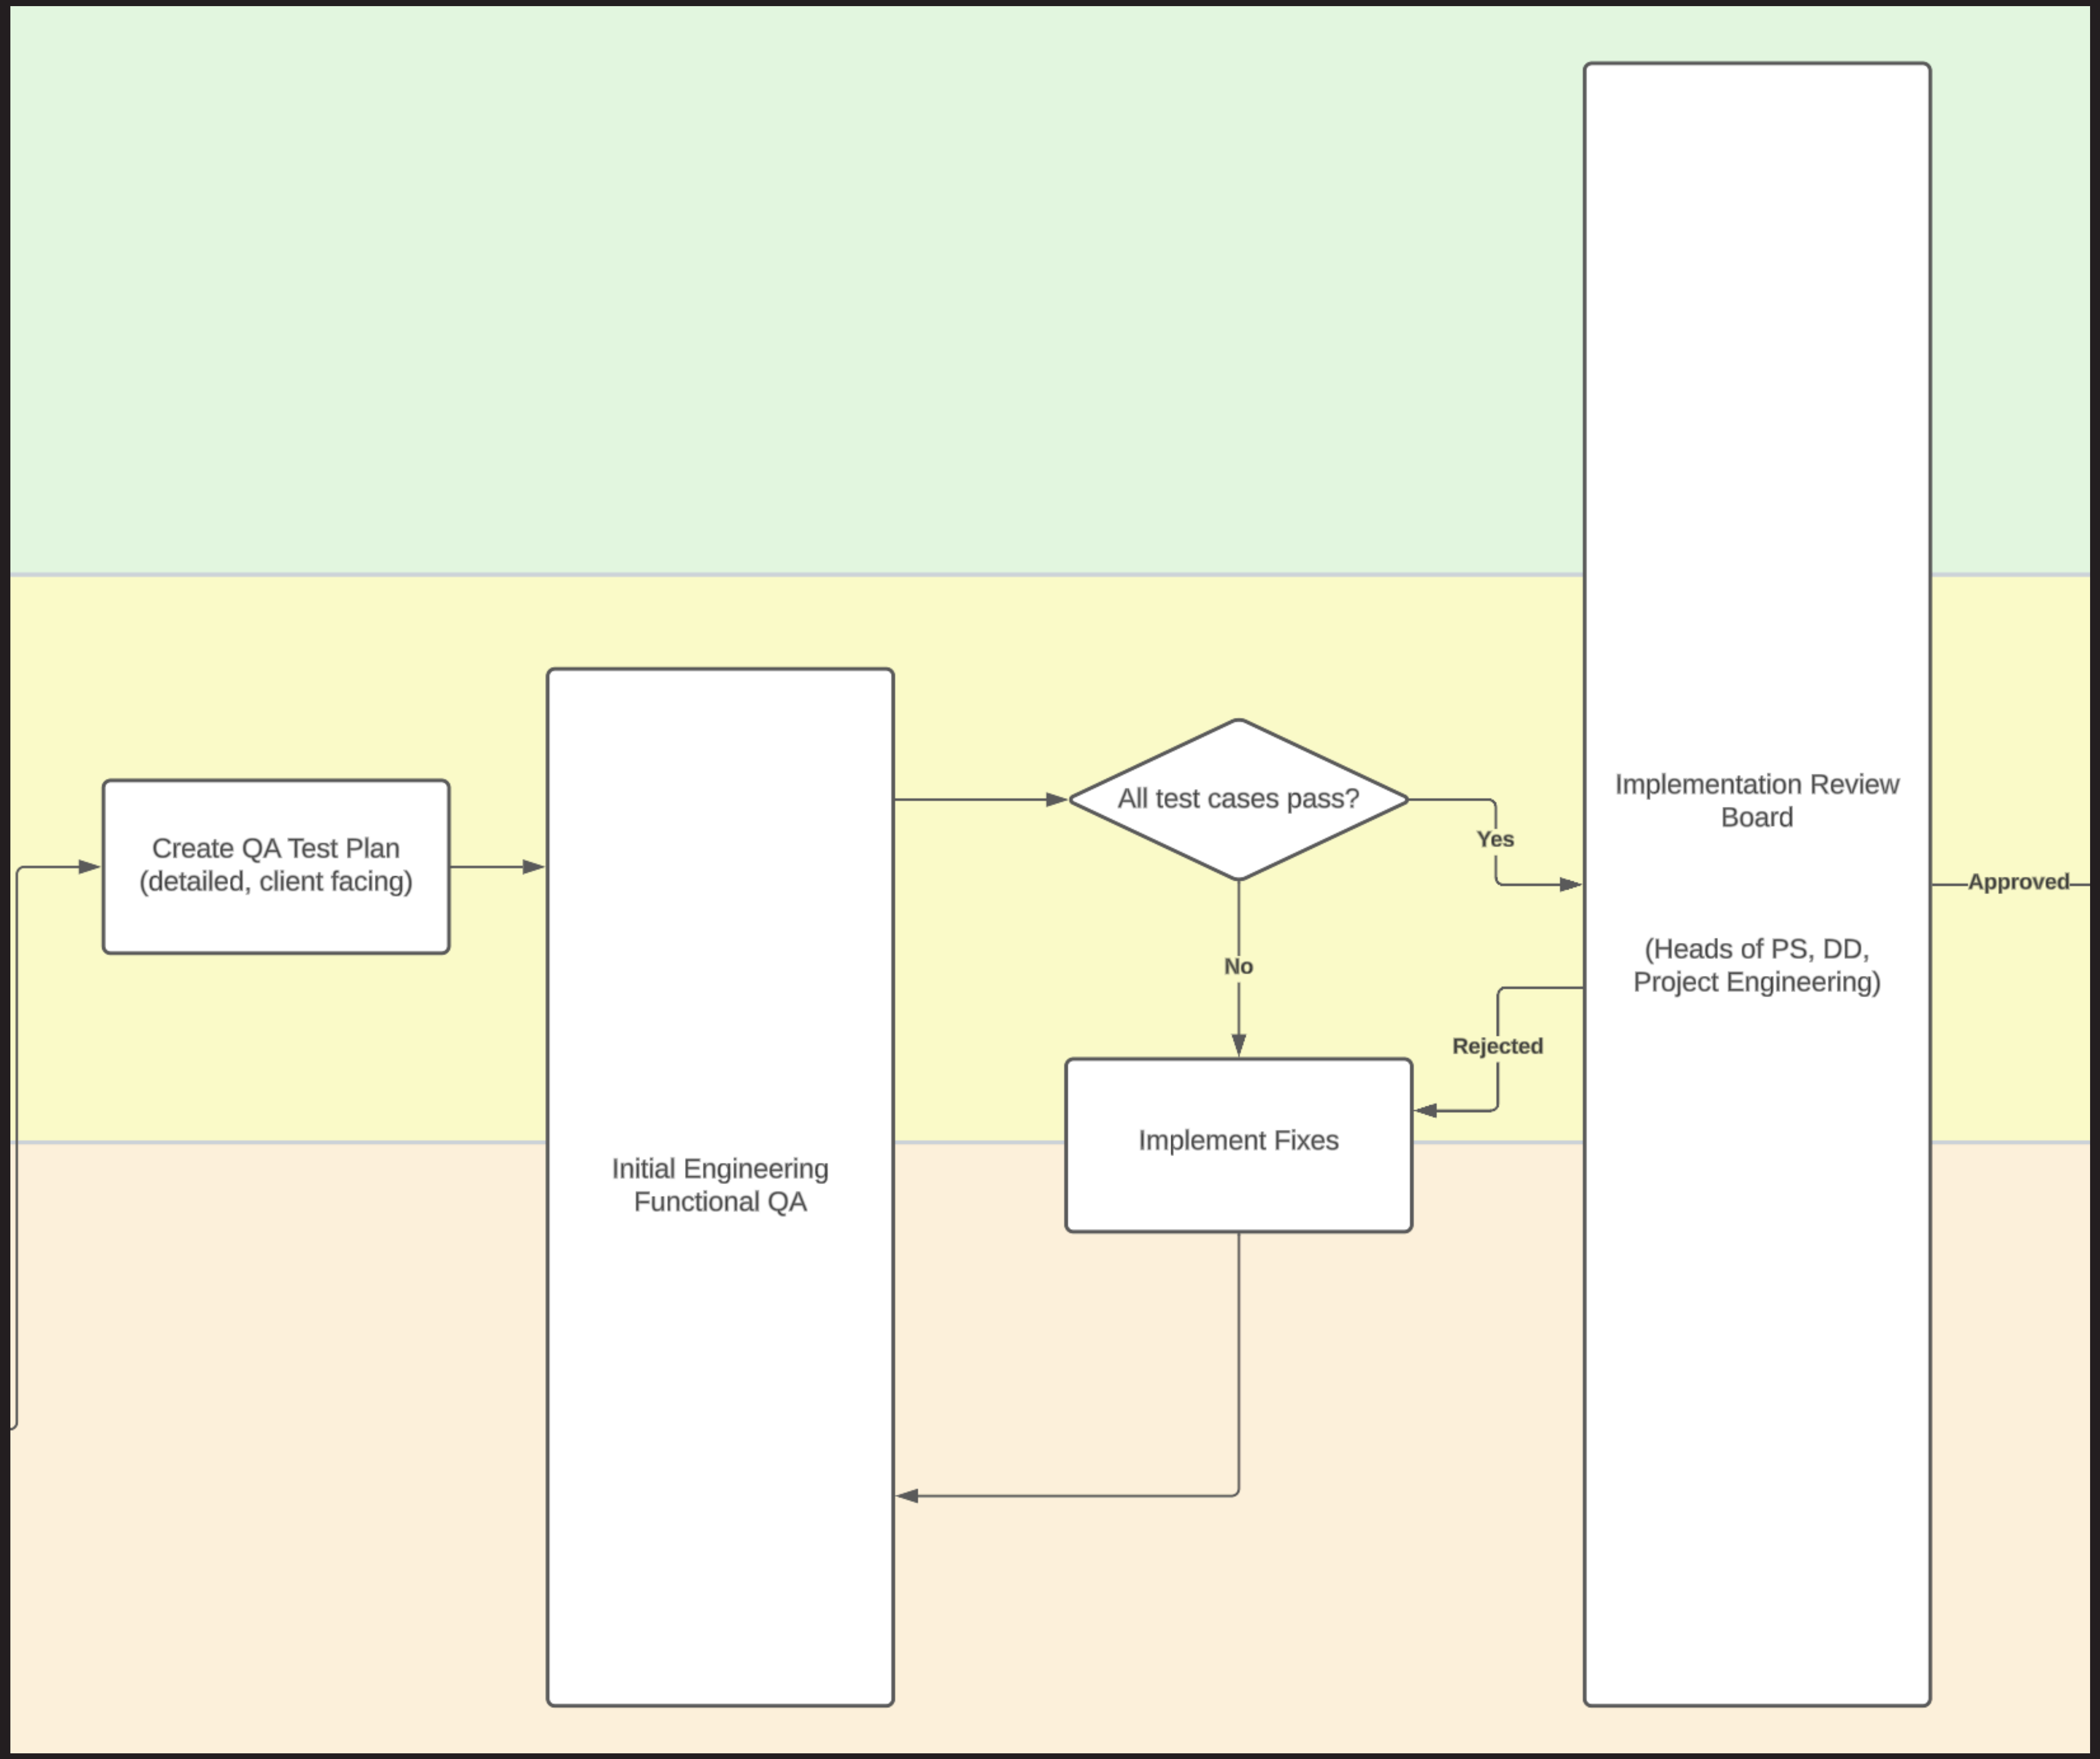

Setting Up QA

The Agent Designer will create the initial QA sheet. This should be detailed and it should be considered client facing because the UAT sheet will be built from it.

It should cover every reasonable scenario, including:

- Happy paths
- API integrations, that should be mocked if not possible early in QA
- Unhappy paths and failures
- Standard behaviour such as 'speak to' and 'who are you'
- That all metrics are being written correctly

There is a template available in Agent Design Notion. A good model to check is that of Pradco / Moultrie Mobile.

Be Ready for Implementation Review

The implementation review board is a new process that will check the agent quality / completeness of QA / metrics and more. See page X for more detail. If the project is rejected at this stage, it may result in delaying UAT.